

# Bayesian Methods for Multimedia Signal Processing

A. Taylan Cemgil

Signal Processing and Communications Lab.



ACM Multimedia 2007, Augsburg, Germany  
Tutorial1a  
September 24, 2007

## Goals of this Tutorial

- Provide a basic understanding of underlying principles of probabilistic modeling and Bayesian inference
- Orientation in the broad literature of Bayesian machine learning and statistical signal processing
- Focus on fundamental concepts rather than technical details,  
... we avoid heavy use of algebra by a graphical notation  
... but there will be some maths

## Goals of this Tutorial

- Model based approach  
... rather than description of algorithms for solving specific problems
- Illustrate with examples how certain problems in multimedia signal analysis can be approached using generic tools
- Motivate participants to investigate further  
... provide alternative perspective to existing solutions  
... and hopefully provide new inspiration

## First Part, Basic Concepts

- Introduction
  - Bayes' Theorem,
  - Trivial toy example to clarify notation
- Graphical Models
  - Bayesian Networks
  - Undirected Graphical models, Markov Random Fields
  - Factor graphs
- Maximum Likelihood, Penalised Likelihood, Bayesian Learning
- Basic Building Blocks in model construction
  - Probability distributions, Exponential family

## Second Part, Models and Applications

- Hidden Markov Models,
  - Tempo tracking, Score-performance matching
  - Inference in Hidden Markov Models
    - \* Forward Backward Algorithm
    - \* Viterbi
    - \* Exact inference by message passing: Belief Propagation
- Linear Dynamical systems, Kalman Filter Models
  - Tracking
  - Computer Accompaniment
  - Kalman Filtering and Smoothing
  - Audio Restoration and Interpolation

- Switching State Space models, Changepoint Models
  - Pitch tracking
  - Particle Filtering
- Nonlinear Dynamical Systems
  - Object tracking in video
  - Particle Filtering, Sequential Monte Carlo
- Markov Random Fields
  - Denoising, Source Separation
  - Markov Chain Monte Carlo, Gibbs sampler
  - Variational Bayes

- Topic-Term Models
  - Latent Semantic indexing
  - Generative aspect model, Latent Dirichlet allocation
- Factorial Models, Sparsity, Model selection
  - Audio Source Separation
  - Polyphonic Pitch Tracking
  - Approximate Inference in Factorial Models
- Final Remarks and Bibliography

## Bayes' Theorem [13, 15]



Thomas Bayes (1702-1761)

What you know about a parameter  $\lambda$  after the data  $\mathcal{D}$  arrive is what you knew before about  $\lambda$  and what the data  $\mathcal{D}$  told you.

$$p(\lambda|\mathcal{D}) = \frac{p(\mathcal{D}|\lambda)p(\lambda)}{p(\mathcal{D})}$$
$$\text{Posterior} = \frac{\text{Likelihood} \times \text{Prior}}{\text{Evidence}}$$

## An application of Bayes' Theorem: "Source Separation"

Given two fair dice with outcomes  $\lambda$  and  $y$ ,

$$\mathcal{D} = \lambda + y$$

What is  $\lambda$  when  $\mathcal{D} = 9$  ?

## An application of Bayes' Theorem: "Source Separation"

$$\mathcal{D} = \lambda + y = 9$$

$\mathcal{D} = \lambda + y$	$y = 1$	$y = 2$	$y = 3$	$y = 4$	$y = 5$	$y = 6$
$\lambda = 1$	2	3	4	5	6	7
$\lambda = 2$	3	4	5	6	7	8
$\lambda = 3$	4	5	6	7	8	<b>9</b>
$\lambda = 4$	5	6	7	8	<b>9</b>	10
$\lambda = 5$	6	7	8	<b>9</b>	10	11
$\lambda = 6$	7	8	<b>9</b>	10	11	12

Bayes theorem "upgrades"  $p(\lambda)$  into  $p(\lambda|\mathcal{D})$ .

But you have to provide an observation model:  $p(\mathcal{D}|\lambda)$

### "Bureaucratical" derivation

Formally we write

$$p(\lambda) = \mathcal{C}(\lambda; [1/6 \ 1/6 \ 1/6 \ 1/6 \ 1/6 \ 1/6])$$

$$p(y) = \mathcal{C}(y; [1/6 \ 1/6 \ 1/6 \ 1/6 \ 1/6 \ 1/6])$$

$$p(\mathcal{D}|\lambda, y) = \delta(\mathcal{D} - (\lambda + y))$$

$$p(\lambda, y|\mathcal{D}) = \frac{1}{p(\mathcal{D})} \times p(\mathcal{D}|\lambda, y) \times p(y)p(\lambda)$$

$$\text{Posterior} = \frac{1}{\text{Evidence}} \times \text{Likelihood} \times \text{Prior}$$

Kronecker delta function denoting a degenerate (deterministic) distribution  $\delta(x) = \begin{cases} 1 & x = 0 \\ 0 & x \neq 0 \end{cases}$

### Prior

$$p(y)p(\lambda)$$

$p(y) \times p(\lambda)$	$y = 1$	$y = 2$	$y = 3$	$y = 4$	$y = 5$	$y = 6$
$\lambda = 1$	1/36	1/36	1/36	1/36	1/36	1/36
$\lambda = 2$	1/36	1/36	1/36	1/36	1/36	1/36
$\lambda = 3$	1/36	1/36	1/36	1/36	1/36	1/36
$\lambda = 4$	1/36	1/36	1/36	1/36	1/36	1/36
$\lambda = 5$	1/36	1/36	1/36	1/36	1/36	1/36
$\lambda = 6$	1/36	1/36	1/36	1/36	1/36	1/36

- A table with indices  $\lambda$  and  $y$
- Each cell denotes the probability  $p(\lambda, y)$

## Likelihood

$$p(\mathcal{D} = 9|\lambda, y)$$

$p(\mathcal{D} = 9 \lambda, y)$	$y = 1$	$y = 2$	$y = 3$	$y = 4$	$y = 5$	$y = 6$
$\lambda = 1$	0	0	0	0	0	0
$\lambda = 2$	0	0	0	0	0	0
$\lambda = 3$	0	0	0	0	0	<b>1</b>
$\lambda = 4$	0	0	0	0	<b>1</b>	0
$\lambda = 5$	0	0	0	<b>1</b>	0	0
$\lambda = 6$	0	0	<b>1</b>	0	0	0

- A table with indices  $\lambda$  and  $y$
- The likelihood is **not** a probability distribution, but a positive function.

## Likelihood $\times$ Prior

$$\phi_{\mathcal{D}}(\lambda, y) = p(\mathcal{D} = 9|\lambda, y)p(\lambda)p(y)$$

$p(\mathcal{D} = 9 \lambda, y)$	$y = 1$	$y = 2$	$y = 3$	$y = 4$	$y = 5$	$y = 6$
$\lambda = 1$	0	0	0	0	0	0
$\lambda = 2$	0	0	0	0	0	0
$\lambda = 3$	0	0	0	0	0	<b>1/36</b>
$\lambda = 4$	0	0	0	0	<b>1/36</b>	0
$\lambda = 5$	0	0	0	<b>1/36</b>	0	0
$\lambda = 6$	0	0	<b>1/36</b>	0	0	0

## Evidence

$$\begin{aligned}
 p(\mathcal{D} = 9) &= \sum_{\lambda, y} p(\mathcal{D} = 9|\lambda, y)p(\lambda)p(y) \\
 &= 0 + 0 + \dots + 1/36 + 1/36 + 1/36 + 1/36 + 0 + \dots + 0 \\
 &= 1/9
 \end{aligned}$$

$p(\mathcal{D} = 9 \lambda, y)$	$y = 1$	$y = 2$	$y = 3$	$y = 4$	$y = 5$	$y = 6$
$\lambda = 1$	0	0	0	0	0	0
$\lambda = 2$	0	0	0	0	0	0
$\lambda = 3$	0	0	0	0	0	<b>1/36</b>
$\lambda = 4$	0	0	0	0	<b>1/36</b>	0
$\lambda = 5$	0	0	0	<b>1/36</b>	0	0
$\lambda = 6$	0	0	<b>1/36</b>	0	0	0

## Posterior

$$p(\lambda, y|\mathcal{D} = 9) = \frac{1}{p(\mathcal{D})}p(\mathcal{D} = 9|\lambda, y)p(\lambda)p(y)$$

$p(\mathcal{D} = 9 \lambda, y)$	$y = 1$	$y = 2$	$y = 3$	$y = 4$	$y = 5$	$y = 6$
$\lambda = 1$	0	0	0	0	0	0
$\lambda = 2$	0	0	0	0	0	0
$\lambda = 3$	0	0	0	0	0	<b>1/4</b>
$\lambda = 4$	0	0	0	0	<b>1/4</b>	0
$\lambda = 5$	0	0	0	<b>1/4</b>	0	0
$\lambda = 6$	0	0	<b>1/4</b>	0	0	0

$$1/4 = (1/36)/(1/9)$$

## Marginal Posterior

$$p(\lambda|\mathcal{D}) = \sum_y \frac{1}{p(\mathcal{D})} p(\mathcal{D}|\lambda, y) p(\lambda) p(y)$$

	$p(\lambda \mathcal{D}=9)$	$y=1$	$y=2$	$y=3$	$y=4$	$y=5$	$y=6$
$\lambda=1$	0	0	0	0	0	0	0
$\lambda=2$	0	0	0	0	0	0	0
$\lambda=3$	<b>1/4</b>	0	0	0	0	0	1/4
$\lambda=4$	<b>1/4</b>	0	0	0	0	1/4	0
$\lambda=5$	<b>1/4</b>	0	0	0	1/4	0	0
$\lambda=6$	<b>1/4</b>	0	0	1/4	0	0	0

## The “proportional to” notation

$$p(\lambda|\mathcal{D}=9) \propto p(\lambda, \mathcal{D}=9) = \sum_y p(\mathcal{D}=9|\lambda, y) p(\lambda) p(y)$$

	$p(\lambda, \mathcal{D}=9)$	$y=1$	$y=2$	$y=3$	$y=4$	$y=5$	$y=6$
$\lambda=1$	0	0	0	0	0	0	0
$\lambda=2$	0	0	0	0	0	0	0
$\lambda=3$	1/36	0	0	0	0	0	<b>1/36</b>
$\lambda=4$	1/36	0	0	0	0	<b>1/36</b>	0
$\lambda=5$	1/36	0	0	0	<b>1/36</b>	0	0
$\lambda=6$	1/36	0	0	<b>1/36</b>	0	0	0

## Exercise

$p(x_1, x_2)$	$x_2=1$	$x_2=2$
$x_1=1$	0.3	0.3
$x_1=2$	0.1	0.3

1. Find the following quantities

- **Marginals:**  $p(x_1)$ ,  $p(x_2)$
- **Conditionals:**  $p(x_1|x_2)$ ,  $p(x_2|x_1)$
- **Posterior:**  $p(x_1, x_2=2)$ ,  $p(x_1|x_2=2)$
- **Evidence:**  $p(x_2=2)$
- $p(\{\})$
- **Max:**  $p(x_1^*) = \max_{x_1} p(x_1|x_2=1)$
- **Mode:**  $x_1^* = \arg \max_{x_1} p(x_1|x_2=1)$
- **Max-marginal:**  $\max_{x_1} p(x_1, x_2)$

2. Are  $x_1$  and  $x_2$  independent ? (i.e., Is  $p(x_1, x_2) = p(x_1)p(x_2)$  ?)

## Answers

$p(x_1, x_2)$	$x_2=1$	$x_2=2$
$x_1=1$	0.3	0.3
$x_1=2$	0.1	0.3

• **Marginals:**

$p(x_1)$		
$x_1=1$	0.6	
$x_1=2$	0.4	

$p(x_2)$	$x_2=1$	$x_2=2$
	0.4	0.6

• **Conditionals:**

$p(x_1 x_2)$	$x_2=1$	$x_2=2$
$x_1=1$	0.75	0.5
$x_1=2$	0.25	0.5

$p(x_2 x_1)$	$x_2=1$	$x_2=2$
$x_1=1$	0.5	0.5
$x_1=2$	0.25	0.75

## Answers

$p(x_1, x_2)$	$x_2 = 1$	$x_2 = 2$
$x_1 = 1$	0.3	0.3
$x_1 = 2$	0.1	0.3

- Posterior:

$p(x_1, x_2 = 2)$	$x_2 = 2$	$p(x_1 x_2 = 2)$	$x_2 = 2$
$x_1 = 1$	0.3	$x_1 = 1$	0.5
$x_1 = 2$	0.3	$x_1 = 2$	0.5

- Evidence:

$$p(x_2 = 2) = \sum_{x_1} p(x_1, x_2 = 2) = 0.6$$

- Normalisation constant:

$$p(\{\}) = \sum_{x_1} \sum_{x_2} p(x_1, x_2) = 1$$

## Answers

$p(x_1, x_2)$	$x_2 = 1$	$x_2 = 2$
$x_1 = 1$	0.3	0.3
$x_1 = 2$	0.1	0.3

- Max: (get the value)

$$\max_{x_1} p(x_1|x_2 = 1) = 0.75$$

- Mode: (get the index)

$$\operatorname{argmax}_{x_1} p(x_1|x_2 = 1) = 1$$

- Max-marginal: (get the “skyline”)  $\max_{x_1} p(x_1, x_2)$

$\max_{x_1} p(x_1, x_2)$	$x_2 = 1$	$x_2 = 2$
	0.3	0.3

## Another application of Bayes’ Theorem: “Model Selection”

Given an unknown number of fair dice with outcomes  $\lambda_1, \lambda_2, \dots, \lambda_n$ ,

$$\mathcal{D} = \sum_{i=1}^n \lambda_i$$

How many dice are there when  $\mathcal{D} = 9$  ?

Assume that any number  $n$  is equally likely

## Another application of Bayes’ Theorem: “Model Selection”

Given all  $n$  are equally likely (i.e.,  $p(n)$  is flat), we calculate (formally)

$$p(n|\mathcal{D} = 9) = \frac{p(\mathcal{D} = 9|n)p(n)}{p(\mathcal{D})} \propto p(\mathcal{D} = 9|n)$$


$$p(\mathcal{D}|n = 1) = \sum_{\lambda_1} p(\mathcal{D}|\lambda_1)p(\lambda_1)$$

$$p(\mathcal{D}|n = 2) = \sum_{\lambda_1} \sum_{\lambda_2} p(\mathcal{D}|\lambda_1, \lambda_2)p(\lambda_1)p(\lambda_2)$$


...

$$p(\mathcal{D}|n = n') = \sum_{\lambda_1, \dots, \lambda_{n'}} p(\mathcal{D}|\lambda_1, \dots, \lambda_{n'}) \prod_{i=1}^{n'} p(\lambda_i)$$

$$p(\mathcal{D}|n) = \sum_{\lambda} p(\mathcal{D}|\lambda, n) p(\lambda|n)$$



## Another application of Bayes' Theorem: "Model Selection"



- Complex models are more flexible but they spread their probability mass
- Bayesian inference inherently prefers "simpler models" – Occam's razor
- Computational burden: We need to sum over all parameters  $\lambda$

## Probabilistic Inference

A huge spectrum of applications – all boil down to computation of

- **expectations** of functions under probability distributions: **Integration**

$$\langle f(x) \rangle = \int_{\mathcal{X}} dx p(x) f(x) \quad \langle f(x) \rangle = \sum_{x \in \mathcal{X}} p(x) f(x)$$

- **modes** of functions under probability distributions: **Optimization**

$$x^* = \operatorname{argmax}_{x \in \mathcal{X}} p(x) f(x)$$

- any "mix" of the above: e.g.,

$$x^* = \operatorname{argmax}_{x \in \mathcal{X}} p(x) = \operatorname{argmax}_{x \in \mathcal{X}} \int_{\mathcal{Z}} dz p(z) p(x|z)$$

## Divide and Conquer

Probabilistic modelling provides a methodology that puts a clear division between

- What to solve : Model Construction
  - Both an Art and Science
  - Highly domain specific
- How to solve : Inference Algorithm
  - Mechanical (In theory! not in practice)
  - Generic

## Applications of Probability Models

- Classification
- Optimal Decision, given a loss function
- Finding interesting (hidden) structure
  - Clustering, Segmentation
  - Dimensionality Reduction
  - Outlier Detection
- Finding a compact representation = Data Compression
- Prediction

## Probability Models

+

## Inference Algorithms

=

## Bayesian Numerical Methods

## Graphical Models

- formal languages for specification of probability models and associated inference algorithms
- historically, introduced in probabilistic expert systems (Pearl 1988) as a visual guide for representing expert knowledge
- today, a standard tool in machine learning, statistics and signal processing

## Graphical Models

- provide graph based algorithms for derivations and computation
- pedagogical insight/motivation for model/algorithm construction
  - Statistics:  
“Kalman filter models and hidden Markov models (HMM) are equivalent upto parametrisation”
  - Signal processing:  
“Fast Fourier transform is an instance of sum-product algorithm on a factor graph”
  - Computer Science:  
“Backtracking in Prolog is equivalent to inference in Bayesian networks with deterministic tables”
- Automated tools for code generation start to emerge, making the design/implement/test cycle shorter




## Important types of Graphical Models

- Useful for Model Construction
  - **Directed Acyclic Graphs (DAG), Bayesian Networks**
  - **Undirected Graphs, Markov Networks, Random Fields**
  - Influence diagrams
  - ...
- Useful for Inference
  - **Factor Graphs**
  - Junction/Clique graphs
  - Region graphs
  - ...


## Directed Graphical models (DAG)

### DAG Example: Two dice



$$p(\mathcal{D}, \lambda, y) = p(\mathcal{D}|\lambda, y)p(\lambda)p(y)$$

### DAG with observations



$$\phi_{\mathcal{D}}(\lambda, y) = p(\mathcal{D} = 9|\lambda, y)p(\lambda)p(y)$$

## Directed Graphical models

- Each random variable is associated with a node in the graph,
- We draw an arrow from  $A \rightarrow B$  if  $p(B | \dots, A, \dots)$  ( $A \in \text{parent}(B)$ ),
- The edges tell us *qualitatively* about the factorization of the joint probability
- For  $N$  random variables  $x_1, \dots, x_N$ , the distribution admits

$$p(x_1, \dots, x_N) = \prod_{i=1}^N p(x_i | \text{parent}(x_i))$$

- Describes in a compact way an algorithm to “generate” the data – “Generative models”

## Undirected Graphical Models

## Examples

Model	Structure	factorization
Full		$p(x_1)p(x_2 x_1)p(x_3 x_1, x_2)p(x_4 x_1, x_2, x_3)$
Markov(2)		$p(x_1)p(x_2 x_1)p(x_3 x_1, x_2)p(x_4 x_2, x_3)$
Markov(1)		$p(x_1)p(x_2 x_1)p(x_3 x_2)p(x_4 x_3)$
		$p(x_1)p(x_2 x_1)p(x_3 x_1)p(x_4)$
Factorized		$p(x_1)p(x_2)p(x_3)p(x_4)$

Removing edges eliminates a term from the conditional probability factors.



## Undirected Graphical Models

- Define a distribution by non-negative *local compatibility functions*  $\phi(x_\alpha)$

$$p(\mathbf{x}) = \frac{1}{Z} \prod_{\alpha} \phi(x_\alpha)$$


where  $\alpha$  runs over **cliques** : fully connected subsets

- Examples



$$p(\mathbf{x}) = \frac{1}{Z} \phi(x_1, x_2) \phi(x_1, x_3) \phi(x_2, x_4) \phi(x_3, x_4) \quad p(\mathbf{x}) = \frac{1}{Z} \phi(x_1, x_2, x_3) \phi(x_2, x_3, x_4)$$


## Possible Model Topologies



## Factor graphs

### Factor graphs [14]

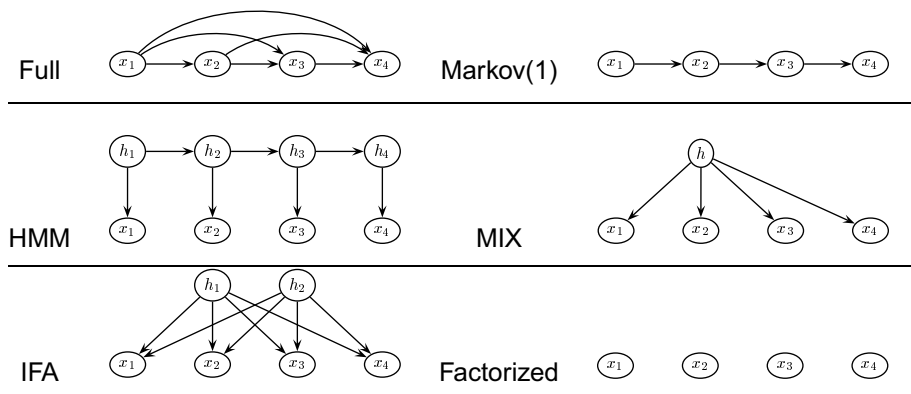
- A bipartite graph. A powerful graphical representation of the inference problem
  - **Factor nodes:** Black squares. Factor potentials (local functions) defining the posterior.
  - **Variable nodes:** White Nodes. Define collections of random variables
  - **Edges:** denote membership. A variable node is connected to a factor node if a member variable is an argument of the local function.



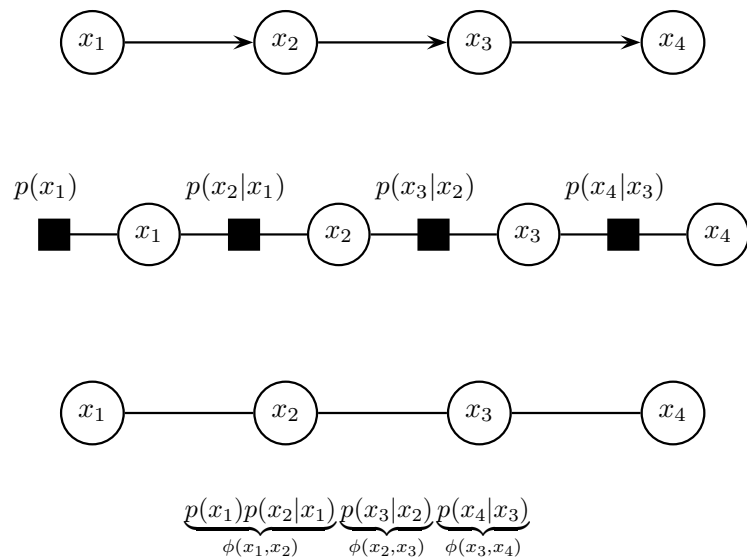
$$\phi_{\mathcal{D}}(\lambda, y) = p(\mathcal{D} = 9 | \lambda, y) p(\lambda) p(y) = \phi_1(\lambda, y) \phi_2(\lambda) \phi_3(y)$$

### Exercise


- For the following Graphical models, write down the factors of the joint distribution and plot an equivalent factor graph and an undirected graph.



### Answer (Markov(1))




### Answer (IFA – Factorial)




$$p(h_1)p(h_2) \prod_{i=1}^4 p(x_i|h_1, h_2)$$

### Answer (IFA – Factorial)



- We can also cluster nodes together



### Inference and Learning

- Data set

$$\mathcal{D} = \{x_1, \dots, x_N\}$$

- Model with parameter  $\lambda$

$$p(\mathcal{D}|\lambda)$$

- Maximum Likelihood (ML)

$$\lambda^{\text{ML}} = \arg \max_{\lambda} \log p(\mathcal{D}|\lambda)$$

- Predictive distribution

$$p(x_{N+1}|\mathcal{D}) \approx p(x_{N+1}|\lambda^{\text{ML}})$$

## Regularisation

- Prior

$$p(\lambda)$$

- Maximum a-posteriori (MAP) : Regularised Maximum Likelihood

$$\lambda^{\text{MAP}} = \arg \max_{\lambda} \log p(\mathcal{D}|\lambda)p(\lambda)$$

- Predictive distribution

$$p(x_{N+1}|\mathcal{D}) \approx p(x_{N+1}|\lambda^{\text{MAP}})$$


## Bayesian Learning

- We treat parameters on the same footing as all other variables
- We integrate over unknown parameters rather than using point estimates (remember the many-dice example)
  - Avoids overfitting
  - Natural setup for online adaptation
  - Model selection
    - (arguably) many problems in music processing are model selection problems

## Bayesian Learning

- Predictive distribution

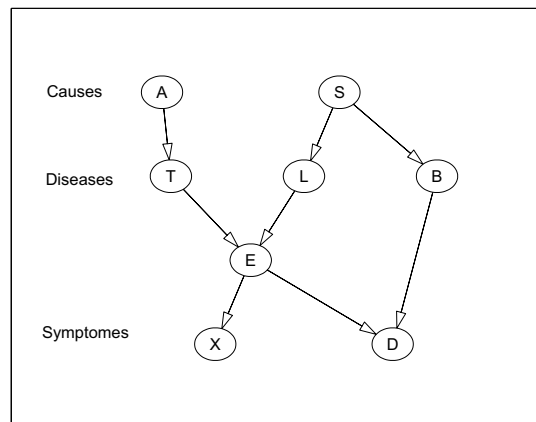
$$p(x_{N+1}|\mathcal{D}) = \int d\lambda \ p(x_{N+1}|\lambda)p(\lambda|\mathcal{D})$$




- Bayesian learning is just inference ...

## Example Applications and Models


## Medical Expert Systems



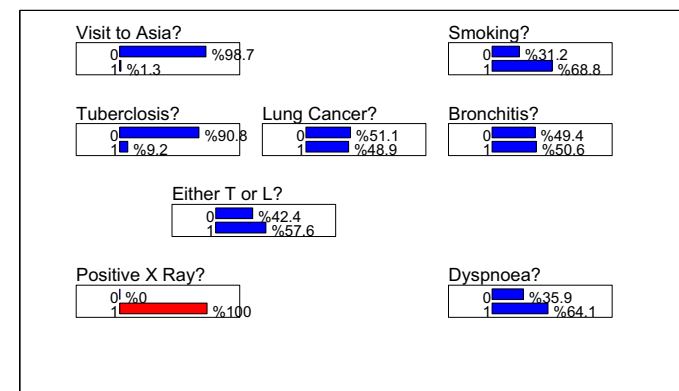
## Medical Expert Systems



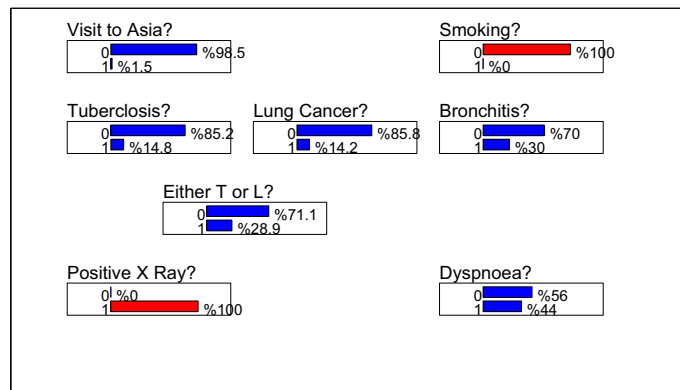
## Medical Expert Systems



## Medical Expert Systems




## Medical Expert Systems




## Model Selection: Variable selection in Polynomial Regression

- Given  $\mathcal{D} = \{t_j, x(t_j)\}_{j=1 \dots J}$ , what is the order  $N$  of the polynomial?

$$x(t) = \sum_{i=0}^N s_{i+1} t^i + \epsilon(t)$$





## Bayesian Variable Selection



- Generalized Linear Model – Column's of  $C$  are the basis vectors
- The exact posterior is a mixture of  $2^W$  Gaussians
- When  $W$  is large, computation of posterior features becomes intractable.

## Regression




All on


Configurations

All off

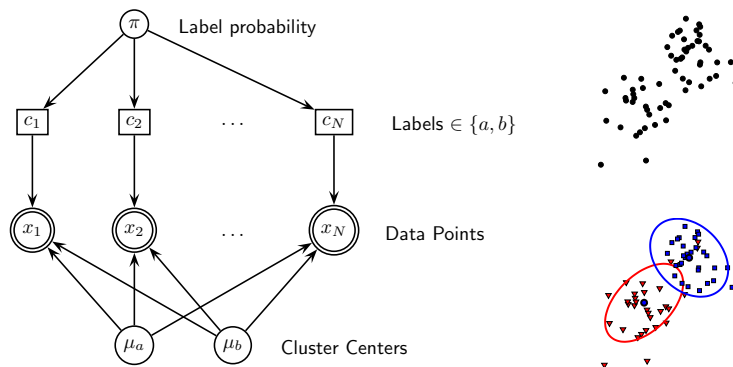
## Regression



## Clustering




## Clustering



$$(\mu_a^*, \mu_b^*, \pi^*) = \underset{\mu_a, \mu_b, \pi}{\operatorname{argmax}} \sum_{c_{1:N}} \prod_{i=1}^N p(x_i | \mu_a, \mu_b, c_i) p(c_i | \pi)$$

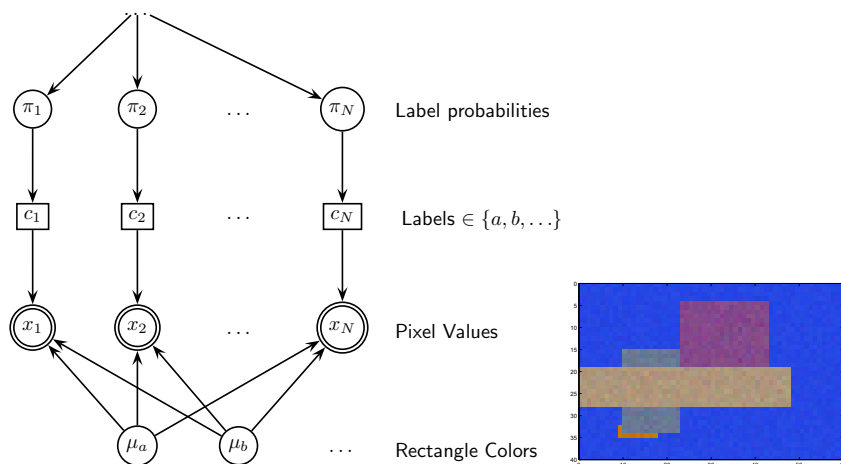
## Computer vision / Cognitive Science

How many rectangles are there in this image?





## Computer vision / Cognitive Science

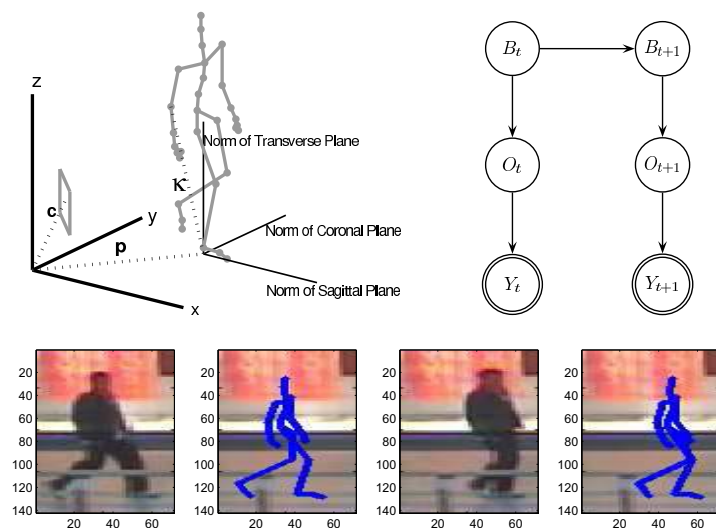


## Computer Vision

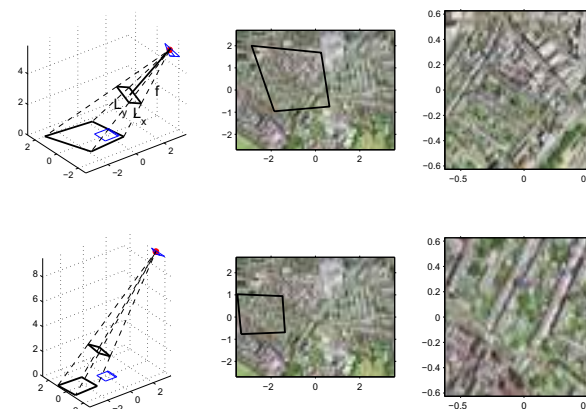
How many people are there in these images?



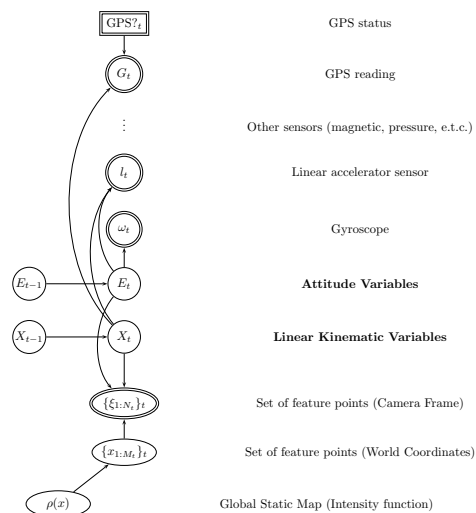
## Visual Tracking



## Navigation, Robotics

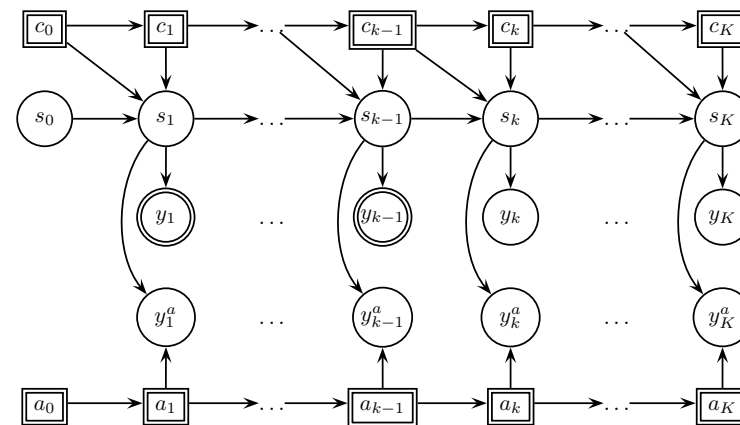


## Navigation, Robotics




## Computer Accompaniment

(Music Plus One, Raphael 2000 [18], Dannenberg and Raphael 2006)



## Audio Restoration



- During download or transmission, some samples of audio are lost
- Estimate missing samples given clean ones



## Examples: Audio Restoration

$$p(x_{-\kappa} | x_{\kappa}) \propto \int d\mathcal{H} p(x_{-\kappa} | \mathcal{H}) p(x_{\kappa} | \mathcal{H}) p(\mathcal{H})$$

$$\mathcal{H} \equiv (\text{parameters, hidden states})$$



## Restoration

(Cemgil and Godsill 2005 [5])

- Piano
  - Signal with missing samples (37%)
  - Reconstruction, 7.68 dB improvement
  - Original
- Trumpet
  - Signal with missing samples (37%)
  - Reconstruction, 7.10 dB improvement
  - Original

## Basic Building Blocks

## Probability Distributions : Exponential Family

- Following distributions are used often as elementary building blocks:
  - Gaussian
  - Gamma, Inverse Gamma, (Exponential, Chi-square, Wishart)
  - Dirichlet
  - Discrete (Categorical), Bernoulli, multinomial
- All of those distributions can be written as

$$p(x|\theta) = \exp\{\theta^\top \psi(x) - A(\theta)\}$$

$$A(\theta) = \log \int_{\mathcal{X}^n} dx \exp(\theta^\top \psi(x)) \quad \text{log-partition function}$$

$\theta$

canonical parameters

$\psi(x)$

sufficient statistics

## Example: Bernoulli

Binary (Bernoulli) random variable  $c = \{0, 1\}$  with probability of success  $w$

$$p(c = 1|w) = w \quad p(c = 0|w) = 1 - w$$

We write

$$\begin{aligned} p(c|w) &= w^c (1 - w)^{1-c} \\ &= \exp(c \log w + (1 - c) \log(1 - w)) \\ &= \exp\left(\log\left(\frac{w}{1 - w}\right)c + \log(1 - w)\right) \\ &= \mathcal{C}(c; w) \end{aligned}$$

$\mathcal{C}$  stays for categorical

### Example, Univariate Gaussian


The Gaussian distribution with mean  $m$  and covariance  $S$  has the form

$$\begin{aligned}
 \mathcal{N}(x; m, S) &= (2\pi S)^{-1/2} \exp\left\{-\frac{1}{2}(x - m)^2/S\right\} \\
 &= \exp\left\{-\frac{1}{2}(x^2 + m^2 - 2xm)/S - \frac{1}{2}\log(2\pi S)\right\} \\
 &= \exp\left\{\frac{m}{S}x - \frac{1}{2S}x^2 - \left(\frac{1}{2}\log(2\pi S) + \frac{1}{2S}m^2\right)\right\} \\
 &= \exp\left\{\underbrace{\begin{pmatrix} m/S \\ -\frac{1}{2S} \end{pmatrix}}_{\theta}^\top \underbrace{\begin{pmatrix} x \\ x^2 \end{pmatrix}}_{\psi(x)} - A(\theta)\right\}
 \end{aligned}$$

Hence by matching coefficients we have

$$\exp\left\{-\frac{1}{2}Kx^2 + hx + g\right\} \Leftrightarrow S = K^{-1} \quad m = K^{-1}h$$

### Example, Gaussian



### Example, Inverse Gamma


The inverse Gamma distribution with shape  $a$  and scale  $b$

$$\begin{aligned}
 \mathcal{IG}(r; a, b) &= \frac{1}{\Gamma(a)} \frac{r^{-(a+1)}}{b^a} \exp\left(-\frac{1}{br}\right) \\
 &= \exp\left(- (a+1) \log r - \frac{1}{br} - \log \Gamma(a) - a \log b\right) \\
 &= \exp\left(\left(\begin{pmatrix} -(a+1) \\ -1/b \end{pmatrix}\right)^\top \begin{pmatrix} \log r \\ 1/r \end{pmatrix} - \log \Gamma(a) - a \log b\right)
 \end{aligned}$$

Hence by matching coefficients, we have

$$\exp\left\{\alpha \log r + \beta \frac{1}{r} + c\right\} \Leftrightarrow a = -\alpha - 1 \quad b = -1/\beta$$

### Example, Inverse Gamma




## Example, Beta

$$\begin{aligned}
 \mathcal{B}(w; a, b) &\equiv \frac{\Gamma(a+b)}{\Gamma(a)\Gamma(b)} w^{a-1} (1-w)^{b-1} \\
 &= \exp((a-1)\log w + (b-1)\log(1-w) - A(a, b)) \\
 &= \exp\left(\begin{pmatrix} a-1 & b-1 \end{pmatrix} \begin{pmatrix} \log w \\ \log(1-w) \end{pmatrix} - A(a, b)\right) \\
 A(a, b) &= \log \Gamma(a) + \log \Gamma(b) - \log \Gamma(a+b)
 \end{aligned}$$

Mean :

$$\langle w \rangle_{\mathcal{B}} = a/(a+b)$$

## Example, Beta



## Conjugate priors: Posterior is in the same family as the prior.

Example: posterior inference for the probability of success  $w$  of a binary (Bernoulli) random variable  $c$

$$\begin{aligned}
 p(c|w) &= \mathcal{C}(c; w) = \exp(c \log w + (1-c) \log(1-w)) \\
 p(w) &= \mathcal{B}(w; a, b)
 \end{aligned}$$

$$\begin{aligned}
 p(w|c) &\propto p(c|w)p(w) \\
 &\propto \exp(c \log w + (1-c) \log(1-w)) \\
 &\quad \times \exp((a-1) \log w + (b-1) \log(1-w)) \\
 &\propto \mathcal{B}(w; a+c, b+(1-c))
 \end{aligned}$$

$$p(w|c) = \begin{cases} \mathcal{B}(w; a+1, b) & c=1 \\ \mathcal{B}(w; a, b+1) & c=0 \end{cases}$$

## Conjugate priors: Posterior is in the same family as the prior.

Example: posterior inference for the variance  $R$  of a zero mean Gaussian.


$$\begin{aligned}
 p(x|R) &= \mathcal{N}(x; 0, R) \\
 p(R) &= \mathcal{IG}(R; a, b)
 \end{aligned}$$

$$\begin{aligned}
 p(R|x) &\propto p(R)p(x|R) \\
 &\propto \exp\left(-(a+1) \log R - (1/b) \frac{1}{R}\right) \exp\left(-(x^2/2) \frac{1}{R} - \frac{1}{2} \log R\right) \\
 &= \exp\left(\begin{pmatrix} -(a+1 + \frac{1}{2}) \\ -(1/b + x^2/2) \end{pmatrix}^\top \begin{pmatrix} \log R \\ 1/R \end{pmatrix}\right) \\
 &\propto \mathcal{IG}(R; a + \frac{1}{2}, \frac{2}{x^2 + 2/b})
 \end{aligned}$$

Like the prior, this is an inverse-Gamma distribution.

## Conjugate priors: Posterior is in the same family as the prior.


Example: posterior inference of variance  $R$  from  $x_1, \dots, x_N$ .




$$\begin{aligned}
 p(R|x) &\propto p(R) \prod_{i=1}^N p(x_i|R) \\
 &\propto \exp\left(-(a+1)\log R - (1/b)\frac{1}{R}\right) \exp\left(-\left(\frac{1}{2}\sum_i x_i^2\right)\frac{1}{R} - \frac{N}{2}\log R\right) \\
 &= \exp\left(\begin{pmatrix} -(a+1+\frac{N}{2}) \\ -(1/b+\frac{1}{2}\sum_i x_i^2) \end{pmatrix}^\top \begin{pmatrix} \log R \\ 1/R \end{pmatrix}\right) \propto \mathcal{IG}(R; a + \frac{N}{2}, \frac{2}{\sum_i x_i^2 + 2/b})
 \end{aligned}$$

Sufficient statistics are **additive**


## Inverse Gamma, $\sum_i x_i^2 = 10$ $N = 10$




## Inverse Gamma, $\sum_i x_i^2 = 100$ $N = 100$



## Inverse Gamma, $\sum_i x_i^2 = 1000$ $N = 1000$



## Example: AR(1) model



$$x_k = Ax_{k-1} + \epsilon_k \quad k = 1 \dots K$$


$\epsilon_k$  is i.i.d., zero mean and normal with variance  $R$ .

### Estimation problem:

Given  $x_0, \dots, x_K$ , determine coefficient  $A$  and variance  $R$  (both scalars).

## AR(1) model, Generative Model notation

$$\begin{aligned} A &\sim \mathcal{N}(A; 0, P) \\ R &\sim \mathcal{IG}(R; \nu, \beta/\nu) \\ x_k | x_{k-1}, A, R &\sim \mathcal{N}(x_k; Ax_{k-1}, R) \quad x_0 = \hat{x}_0 \end{aligned}$$



Observed variables are shown with double circles


## AR(1) Model. Bayesian Posterior Inference

$$p(A, R | x_0, x_1, \dots, x_K) \propto p(x_1, \dots, x_K | x_0, A, R) p(A, R)$$

$$\text{Posterior} \propto \text{Likelihood} \times \text{Prior}$$


Using the Markovian (conditional independence) structure we have

$$p(A, R | x_0, x_1, \dots, x_K) \propto \left( \prod_{k=1}^K p(x_k | x_{k-1}, A, R) \right) p(A) p(R)$$



## Numerical Example

Suppose  $K = 1$ ,



By Bayes' Theorem and the structure of AR(1) model

$$\begin{aligned} p(A, R | x_0, x_1) &\propto p(x_1 | x_0, A, R) p(A) p(R) \\ &= \mathcal{N}(x_1; Ax_0, R) \mathcal{N}(A; 0, P) \mathcal{IG}(R; \nu, \beta/\nu) \end{aligned}$$

## Numerical Example


$$\begin{aligned}
 p(A, R|x_0, x_1) &\propto p(x_1|x_0, A, R)p(A)p(R) \\
 &= \mathcal{N}(x_1; Ax_0, R)\mathcal{N}(A; 0, P)\mathcal{IG}(R; \nu, \beta/\nu) \\
 &\propto \exp\left(-\frac{1}{2}\frac{x_1^2}{R} + x_0x_1\frac{A}{R} - \frac{1}{2}\frac{x_0^2A^2}{R} - \frac{1}{2}\log 2\pi R\right) \\
 &\quad \exp\left(-\frac{1}{2}\frac{A^2}{P}\right) \exp\left(-(\nu+1)\log R - \frac{\nu}{\beta R}\right)
 \end{aligned}$$

This posterior has a nonstandard form

$$\exp\left(\alpha_1\frac{1}{R} + \alpha_2\frac{A}{R} + \alpha_3\frac{A^2}{R} + \alpha_4\log R + \alpha_5A^2\right)$$

## Numerical Example, the prior $p(A, R)$


Equiprobability contour of  $p(A)p(R)$



$$A \sim \mathcal{N}(A; 0, 1.2) \quad R \sim \mathcal{IG}(R; 0.4, 250)$$

$$\text{Suppose: } x_0 = 1 \quad x_1 = -6 \quad x_1 \sim \mathcal{N}(x_1; Ax_0, R)$$

## Numerical Example, the posterior $p(A, R|x)$



Note the bimodal posterior with  $x_0 = 1, x_1 = -6$

- $A \approx -6 \Leftrightarrow$  low noise variance  $R$ .
- $A \approx 0 \Leftrightarrow$  high noise variance  $R$ .

## Remarks

- The point estimates such as ML or MAP are not always representative about the solution
- (Unfortunately), exact posterior inference is only possible for few special cases
- Even very simple models can lead easily to complicated posterior distributions
- Ambiguous data usually leads to a multimodal posterior, each mode corresponding to one possible explanation




## Remarks

- *A-priori* independent variables often become dependent *a-posteriori* (“Explaining away”)
- The difficulty of an inference problem depends, among others, upon the particular “parameter regime” and observed data sequence


## Approximate Inference

### A Toy Model




$$\begin{aligned}
 s_1 &\sim p(s_1) = \mathcal{N}(s_1; \mu_1, P_1) \\
 s_2 &\sim p(s_2) = \mathcal{N}(s_2; \mu_2, P_2) \\
 x|s_1, s_2 &\sim p(x|s_1, s_2) = \mathcal{N}(x; s_1 + s_2, R)
 \end{aligned}$$


### Gibbs Sampling




# Gibbs Sampling




# Gibbs Sampling




# Gibbs Sampling, $t = 20$



# Gibbs Sampling, $t = 100$



## Gibbs Sampling, $t = 250$



## Gibbs Sampling

- A remarkable fact is that we can estimate any desired expectation by ergodic averages

$$\langle f(s) \rangle_{\mathcal{P}} \approx \frac{1}{t - t_0} \sum_{n=t_0}^t f(s^{(n)})$$

- Consecutive samples  $s^{(t)}$  are dependent but we can “pretend” as if they are independent!
- The sequence of samples are obtained from a Markov chain, hence the name MCMC

## Variational Bayes (VB), mean field

We will approximate the posterior  $\mathcal{P}$  with a simpler distribution  $\mathcal{Q}$ .

$$\mathcal{P} = \frac{1}{Z_x} p(x = \hat{x} | s_1, s_2) p(s_1) p(s_2)$$

$$\mathcal{Q} = q(s_1) q(s_2)$$

Here, we choose

$$q(s_1) = \mathcal{N}(s_1; m_1, S_1) \quad q(s_2) = \mathcal{N}(s_2; m_2, S_2)$$

A “measure of fit” between distributions is the KL divergence

## Kullback-Leibler (KL) Divergence

- A “quasi-distance” between two distributions  $\mathcal{P} = p(x)$  and  $\mathcal{Q} = q(x)$ .

$$KL(\mathcal{P} || \mathcal{Q}) \equiv \int_{\mathcal{X}} dx p(x) \log \frac{p(x)}{q(x)} = \langle \log \mathcal{P} \rangle_{\mathcal{P}} - \langle \log \mathcal{Q} \rangle_{\mathcal{P}}$$

- Unlike a metric, (in general) it is not symmetric,

$$KL(\mathcal{P} || \mathcal{Q}) \neq KL(\mathcal{Q} || \mathcal{P})$$

- But it is non-negative (by Jensen’s Inequality)

$$\begin{aligned} KL(\mathcal{P} || \mathcal{Q}) &= - \int_{\mathcal{X}} dx p(x) \log \frac{q(x)}{p(x)} \\ &\geq - \log \int_{\mathcal{X}} dx p(x) \frac{q(x)}{p(x)} = - \log \int_{\mathcal{X}} dx q(x) = - \log 1 = 0 \end{aligned}$$

## OSSS example, cont.

Let the approximating distribution be factorized as

$$\mathcal{Q} = q(s_1)q(s_2)$$

$$q(s_1) = \mathcal{N}(s_1; m_1, S_1) \quad q(s_2) = \mathcal{N}(s_2; m_2, S_2)$$

The  $m_i$  and  $S_j$  are the *variational* parameters to be optimized to minimize

$$KL(\mathcal{Q} || \mathcal{P}) = \langle \log \mathcal{Q} \rangle_{\mathcal{Q}} - \left\langle \log \underbrace{\frac{1}{Z_x} \phi(s_1, s_2)}_{=\mathcal{P}} \right\rangle_{\mathcal{Q}} \quad (1)$$

## The form of the mean field solution

$$\begin{aligned} 0 &\leq \langle \log q(s_1)q(s_2) \rangle_{q(s_1)q(s_2)} + \log Z_x - \langle \log \phi(s_1, s_2) \rangle_{q(s_1)q(s_2)} \\ \log Z_x &\geq \langle \log \phi(s_1, s_2) \rangle_{q(s_1)q(s_2)} - \langle \log q(s_1)q(s_2) \rangle_{q(s_1)q(s_2)} \\ &\equiv -F(p; q) + H(q) \end{aligned} \quad (2)$$

Here,  $F$  is the *energy* and  $H$  is the *entropy*. We need to maximize the right hand side.

$$\text{Evidence} \geq -\text{Energy} + \text{Entropy}$$

Note r.h.s. is a **lower bound** [16]. The mean field equations **monotonically** increase this bound. Good for assessing convergence and debugging computer code.

## The form of the solution


- No direct analytical solution
- We obtain fixed point equations in closed form

$$q(s_1) \propto \exp(\langle \log \phi(s_1, s_2) \rangle_{q(s_2)})$$

$$q(s_2) \propto \exp(\langle \log \phi(s_1, s_2) \rangle_{q(s_1)})$$


Note the nice symmetry

## Variational Message Passing on a Factor Graph



- **Factor nodes:** Factor potentials (local functions) defining the posterior  $\mathcal{P}$ .
- **Variable nodes:** Now, think of them as “factors” of the approximating distribution  $\mathcal{Q}$ . (Caution – non standard interpretation!)


## Fixed Point Iteration



$$\log q(s_1) \leftarrow \log p(s_1) + \langle \log p(x = \hat{x} | s_1, s_2) \rangle_{q(s_2)}$$

$$\log q(s_2) \leftarrow \log p(s_2) + \langle \log p(x = \hat{x} | s_1, s_2) \rangle_{q(s_1)}$$

## VB Convergence



## Direct Link to Expectation-Maximisation (EM) [12]

Suppose we choose one of the distributions degenerate, i.e.

$$\tilde{q}(s_2) = \delta(s_2 - \tilde{m})$$

where  $\tilde{m}$  corresponds to the “location parameter” of  $\tilde{q}(s_2)$ . We need to find the closest degenerate distribution to the actual mean field solution  $q(s_2)$ , hence we take one more KL and minimize

$$\tilde{m} = \underset{\xi}{\operatorname{argmin}} KL(\delta(s_2 - \xi) || q(s_2))$$

It can be shown that this leads exactly to the EM fixed point iterations.

## Iterated Conditional Modes (ICM) [2, 11]

If we choose both distributions degenerate, i.e.

$$\tilde{q}(s_1) = \delta(s_1 - \tilde{m}_1)$$

$$\tilde{q}(s_2) = \delta(s_2 - \tilde{m}_2)$$

It can be shown that this leads exactly to the ICM fixed point iterations. This algorithm is equivalent to coordinate ascent in the original posterior surface  $\phi(s_1, s_2)$ .

$$\tilde{m}_1 = \underset{s_1}{\operatorname{argmax}} \phi(s_1, s_2 = \tilde{m}_2)$$

$$\tilde{m}_2 = \underset{s_2}{\operatorname{argmax}} \phi(s_1 = \tilde{m}_1, s_2)$$

## ICM, EM, VB ...

For OSSS, all algorithms are identical. This is in general not true.

While algorithmic details are very similar, there can be big qualitative differences in terms of fixed points.





Figure 1: Left, ICM, Right VB. EM is similar to ICM in this AR(1) example.

## Models and Applications

## Time series models and Inference, Terminology

In music signal processing and machine learning many phenomena are modelled by dynamical models




- $x$  is the latent state (tempo, pitch, velocity, attitude, class label, ...)
- $y$  are observations (samples, onsets, sensor reading, pixels, features, ...)
- In a full Bayesian setting,  $x$  includes unknown model parameters

## Online Inference, Terminology

### • Filtering: $p(x_k | y_{1:k})$

- Distribution of current state given all past information
- Realtime/Online/Sequential Processing




### • Potentially confusing misnomer:

- More general than “digital filtering” (convolution) in DSP – but algorithmically related for some models (KFM)

## Online Inference, Terminology

- **Prediction**  $p(y_{k:K}, x_{k:K} | y_{1:k-1})$


- evaluation of possible future outcomes; like filtering without observations




- Accompaniment, Tracking, Restoration

## Offline Inference, Terminology

- **Smoothing**  $p(x_{0:K}|y_{1:K})$ ,  
**Most likely trajectory – Viterbi path**  $\arg \max_{x_{0:K}} p(x_{0:K}|y_{1:K})$   
 better estimate of past states, essential for learning




- **Interpolation**  $p(y_k, x_k | y_{1:k-1}, y_{k+1:K})$   
fill in lost observations given past and future



## Hidden Markov Model [17]

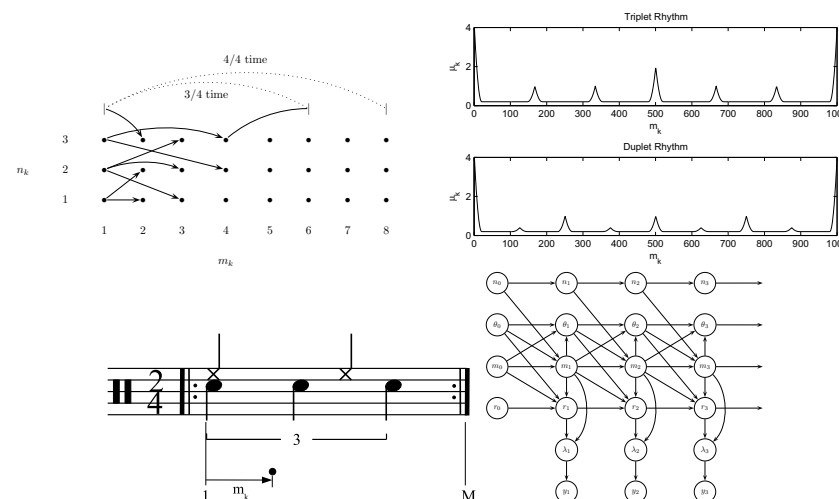
- Mixture model evolving in time



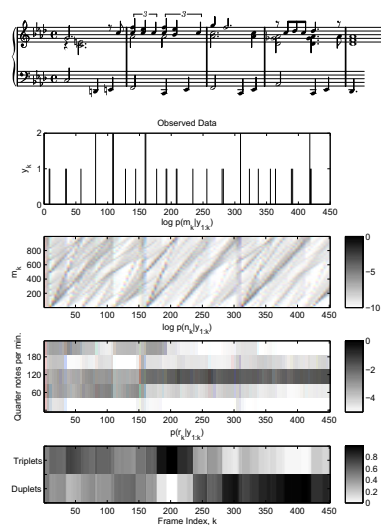
- Observations  $y_k$  are continuous or discrete
- Latent variables  $x_k$  are discrete
  - Represents the fading memory of the process
- Exact inference possible if  $x_k$  has a “small” number of states

## Tempo, Rhythm, Meter analysis

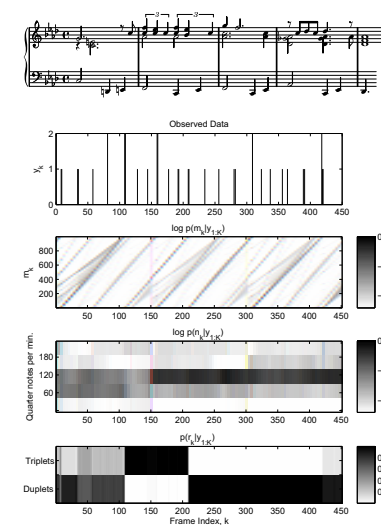
### Bar Pointer Model (Whiteley, Cemgil, Godsill 2006)



## Filtering




## Smoothing

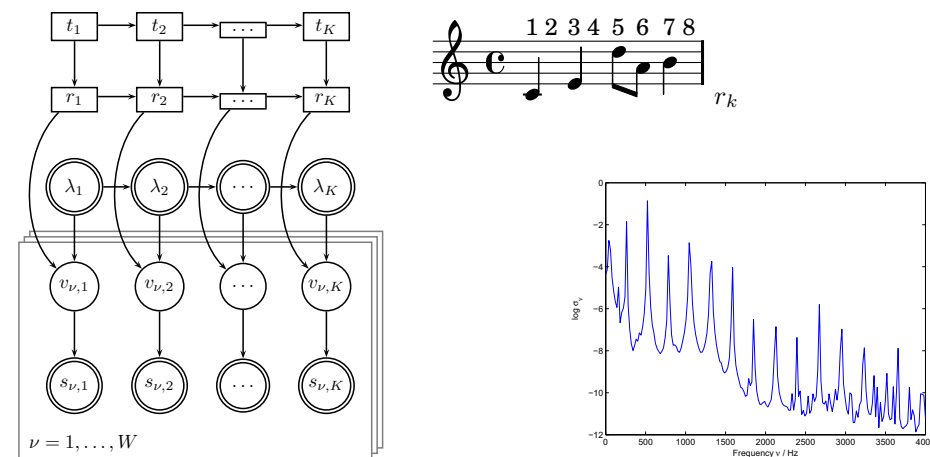


## Score-Performance matching (Peeling, Cemgil, Godsill)

- Given a musical score, associate note events with the audio



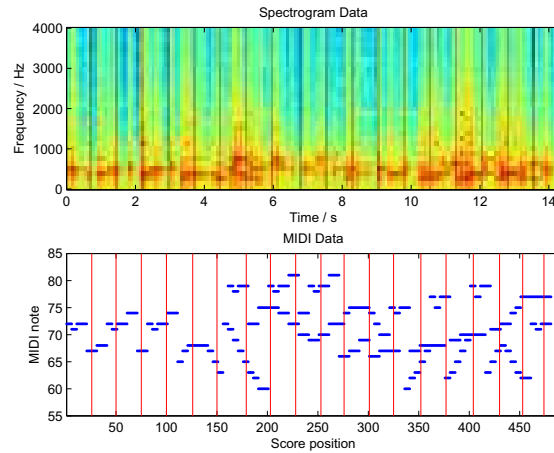
## Score-Performance matching - Graphical Model




$$v_{\nu,\tau} \sim \mathcal{IG}(v_{\nu,\tau}; a, 1/(a\lambda\sigma_\nu(r_\tau)))$$



## Score-Performance matching



## Exact Inference in HMM, Forward/Backward Algorithm




### • Forward Pass

$$\begin{aligned}
 p(y_{1:K}) &= \sum_{x_{1:K}} p(y_{1:K}|x_{1:K})p(x_{1:K}) \\
 &= \underbrace{\sum_{x_K} p(y_K|x_K) \sum_{x_{K-1}} p(x_K|x_{K-1}) \cdots \sum_{x_2} p(x_3|x_2) p(y_2|x_2)}_{\alpha_K} \underbrace{\sum_{x_1} p(x_2|x_1)}_{\alpha_2} \underbrace{p(y_1|x_1) p(x_1)}_{\alpha_1}
 \end{aligned}$$

### • Backward Pass

$$p(y_{1:K}) = \sum_{x_1} p(x_1) p(y_1|x_1) \cdots \underbrace{\sum_{x_{K-1}} p(x_{K-1}|x_{K-2}) p(y_{K-1}|x_{K-1})}_{\beta_{K-2}} \underbrace{\sum_{x_K} p(x_K|x_{K-1}) p(y_K|x_K)}_{\beta_{K-1}} \underbrace{1}_{\beta_K}$$

## Exact Inference in HMM, Viterbi Algorithm



### • Merely replace sum by max, equivalent to dynamic programming

### • Forward Pass

$$\begin{aligned}
 p(y_{1:K}|x_{1:K}^*) &= \max_{x_{1:K}} p(y_{1:K}|x_{1:K})p(x_{1:K}) \\
 &= \underbrace{\max_{x_K} p(y_K|x_K) \max_{x_{K-1}} p(x_K|x_{K-1}) \cdots \max_{x_2} p(x_3|x_2) p(y_2|x_2)}_{\alpha_K} \underbrace{\max_{x_1} p(x_2|x_1)}_{\alpha_2} \underbrace{p(y_1|x_1) p(x_1)}_{\alpha_1}
 \end{aligned}$$

### • Backward Pass

$$p(y_{1:K}|x_{1:K}^*) = \max_{x_1} p(x_1) p(y_1|x_1) \cdots \underbrace{\max_{x_{K-1}} p(x_{K-1}|x_{K-2}) p(y_{K-1}|x_{K-1})}_{\beta_{K-2}} \underbrace{\max_{x_K} p(x_K|x_{K-1}) p(y_K|x_K)}_{\beta_{K-1}} \underbrace{1}_{\beta_K}$$

## Exact Inference on general factor graphs

- When the factor graph is a tree, one can define a local message propagation
  - If factor graph is not a tree, one can always do this by clustering nodes together

### • Sum-product

- Generalises Forward/Backward
- Rule:


“The message sent from a node  $v$  on an edge  $e$  is the product of the local function at  $v$  (or the unit function if  $v$  is a variable node) with all messages received at  $v$  on edges other than  $e$ , summarized for the variable associated with  $e$ .”

### • Max-product

- Generalises Viterbi

Look at the seminal tutorial paper by Kschischang, Frey and Loeliger [14] on factor graphs.

## Exact Inference on general factor graphs



variable to local function:

$$\mu_{x \rightarrow f}(x) = \prod_{h \in n(x) \setminus \{f\}} \mu_{h \rightarrow x}(x)$$

local function to variable:

$$\mu_{f \rightarrow x}(x) = \sum_{\sim x} \left( f(X) \prod_{y \in n(f) \setminus \{x\}} \mu_{y \rightarrow f}(y) \right)$$

## Kalman Filter Models, Linear Dynamical Systems

- The latent variables  $s_k$  and observations  $y_k$  are continuous
- The transition and observations models are linear
  - Example: a point moving on the real line
  - A deterministic dynamical system with two state variables

$$\mathbf{s}_k = \begin{pmatrix} \text{position} \\ \text{velocity} \end{pmatrix}_k = \begin{pmatrix} 1 & 1 \\ 0 & 1 \end{pmatrix} \mathbf{s}_{k-1} = \mathbf{A} \mathbf{s}_{k-1}$$

$$y_k = \text{position}_k = \begin{pmatrix} 1 & 0 \end{pmatrix} \mathbf{s}_k = \mathbf{C} \mathbf{s}_k$$


## Tracking

- We allow random (unknown) accelerations

$$\begin{aligned} \mathbf{s}_k &= \begin{pmatrix} 1 & 1 \\ 0 & 1 \end{pmatrix} \mathbf{s}_{k-1} + \epsilon_k \\ &= \mathbf{A} \mathbf{s}_{k-1} + \epsilon_k \end{aligned}$$

$$\begin{aligned} y_k &= \begin{pmatrix} 1 & 0 \end{pmatrix} \mathbf{s}_k + \nu_k \\ &= \mathbf{C} \mathbf{s}_k + \nu_k \end{aligned}$$

## Tracking




- In generative model notation

$$\mathbf{s}_k \sim \mathcal{N}(\mathbf{s}_k; \mathbf{A} \mathbf{s}_{k-1}, Q)$$

$$y_k \sim \mathcal{N}(y_k; \mathbf{C} \mathbf{s}_k, R)$$

- Tracking = estimating the latent state of the system = Kalman filtering

## Kalman Filtering and Smoothing (two filter formulation)



- Forward Pass



$$p(y_{1:K}) = \underbrace{\int_{x_K} p(y_K|x_K) \int_{x_{K-1}} p(x_K|x_{K-1}) \dots \int_{x_2} p(x_3|x_2) p(y_2|x_2)}_{\alpha_K} \underbrace{\int_{x_1} p(x_2|x_1) p(y_1|x_1)}_{\alpha_1} \underbrace{p(x_1)}_{\alpha_1|0}$$

- Backward Pass



$$p(y_{1:K}) = \int_{x_1} p(x_1) p(y_1|x_1) \dots \underbrace{\int_{x_{K-1}} p(x_{K-1}|x_{K-2}) p(y_{K-1}|x_{K-1})}_{\beta_{K-2}} \underbrace{\int_{x_K} p(x_K|x_{K-1}) p(y_K|x_K)}_{\beta_{K-1}} \underbrace{1}_{\beta_K}$$

- Replace summation by integration



$$p(s_1)$$




$$p(y_1|s_1)p(s_1)$$




$$p(s_2|y_1) \propto \int ds_1 p(s_2|s_1) p(y_1|s_1) p(s_1)$$



$$p(y_2|s_2)p(s_2|y_1)$$




$$p(s_5|y_{1:5})$$



## Computer Accompaniment

(Music Plus One, Raphael 2000 [18], Dannenberg and Raphael 2006)



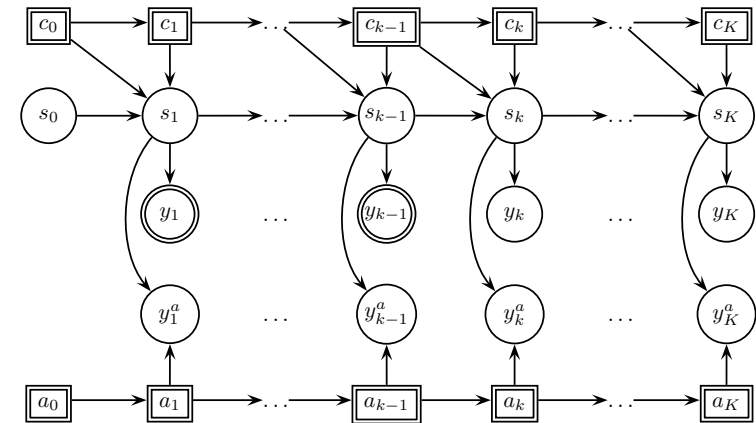
- $c_k$  are score positions of notes of the soloist and  $l_k = c_k - c_{k-1}$

$$\mathbf{s}_k = \begin{pmatrix} 1 & l_k \\ 0 & 1 \end{pmatrix} \mathbf{s}_{k-1} + \epsilon_k = \mathbf{A}_k \mathbf{s}_{k-1} + \epsilon_k \quad y_k = C \mathbf{s}_k + \nu_k$$

$$\epsilon_k \sim \mathcal{N}(\epsilon; 0, Q_k)$$


$$\nu_k \sim \mathcal{N}(\nu; m_k, R_k) \quad (\text{note } k \text{ dependent mean and variance!})$$

## Music Plus One



- Note that this is ruthless simplification, see Chris Raphaels' papers...


## Switching State Space models



- We introduce latent switch variables to switch between different transition and observation models
- Powerful framework for modelling nonstationary processes and nonlinear dynamical systems

## Inference in Switching State Space models


- Unlike HMM's or KFM's, summing over  $c_k$  does not simplify the filtering density.
- Number of Gaussian kernels to represent exact filtering density  $p(c_k, s_k | y_{1:k})$  increases exponentially



- Bad news: exact inference problem is shown to be NP hard


## Example

Suppose that a score can consist of only two notes:




## Sequential Monte Carlo (Particle Filtering)

- Main idea: Select a branch to expand with a probability proportional to the evidence



## Particle Filtering for tracking



## Sequential Monte Carlo

- This variant is known as Mixture Kalman Filter or Rao-Blackwellized Particle filter (Chen and Liu 2001 [9], Cemgil 2002 [6], Hainsworth and MacLeod 2003)
- (For this model) algorithmically similar to Breadth first search/Multi Hypothesis Tracking/Genetic algorithms
- Generic tool for inference with a rich background theory (Doucet, et. al. 2001, Del Moral, "Feynman-Kac Formulae", 2005)
- Many applications in various fields
  - Robotics, Navigation, Econometrics,...

## Changepoint models

$$r_k \sim p(r_k | r_{k-1})$$


Indicators  $\in \{\text{new}, \text{reg}\}$

$$\theta_k \sim [r_k = \text{reg}] \underbrace{f(\theta_k | \theta_{k-1})}_{\text{Transition}} + [r_k = \text{new}] \underbrace{\pi(\theta_k)}_{\text{Reinitialization}}$$


Latent State

$$y_k \sim p(y_k | \theta_k)$$

Observations



## Example: Single Key, Onsets



- Each changepoint denotes the onset of a new audio event

## Dynamic Harmonic Model (Cemgil et. al. 2005, 2006) [4, 7]

$$r_k | r_{k-1} \sim p(r_k | r_{k-1})$$

$$s_k | s_{k-1}, r_k \sim \underbrace{[r_k = 0] \mathcal{N}(A s_{k-1}, Q)}_{\text{reg}} + \underbrace{[r_k = 1] \mathcal{N}(0, S)}_{\text{new}}$$

$$y_k | s_k \sim \mathcal{N}(C s_k, R)$$




$$A = \begin{pmatrix} G_\omega & & & \\ & G_\omega^2 & & \\ & & \ddots & \\ & & & G_\omega^H \end{pmatrix}^N \quad G_\omega = \rho_k \begin{pmatrix} \cos(\omega) & -\sin(\omega) \\ \sin(\omega) & \cos(\omega) \end{pmatrix}$$

damping factor  $0 < \rho_k < 1$ , framelength  $N$  and damped sinusoidal basis matrix  $C$  of size  $N \times 2H$


## Monophonic model [7]

- We introduce a pitch label indicator  $m$
- At each time  $k$ , the process can be in one of the  $\{\text{"mute", "sound"}\} \times M$  states.



## Monophonic Pitch Tracking


Monophonic Pitch Tracking = Online estimation (filtering) of  $p(r_k, m_k | y_{1:k})$ .



- If pitch is constant exact inference is possible


## Tracking Pitch Variations

- Allow  $m$  to change with  $k$ . We take a fine grid Piano-roll, e.g.  $Q = 2^{1/128}$




- Intractable, need to run a particle filter

## Real Data Results



Top: F major scale played on an electric bass.  
Bottom: Estimated MAP configuration  $(r, m)_{1:T}$ .

## Real Data Results



A finer analysis with  $Q = 2^{1/48}$  reveals that the 5'th and 7'th degree of the scale are intonated slightly low.

## Polyphony: Factorial Dynamic Harmonic Model [4]

$$r_{0,\nu} \sim \mathcal{C}(r_{0,\nu}; \pi_{0,\nu})$$

$$\theta_{0,\nu} \sim \mathcal{N}(\theta_{0,\nu}; \mu_\nu, P_\nu)$$

$$r_{k,\nu} | r_{k-1,\nu} \sim \mathcal{C}(r_{k,\nu}; \pi_\nu(r_{k-1,\nu}))$$


Changepoint indicator

$$\theta_{k,\nu} | \theta_{k-1,\nu} \sim \mathcal{N}(\theta_{k,\nu}; A_\nu(r_k) \theta_{k-1,\nu}, Q_\nu(r_k))$$

Latent state

$$y_k | \theta_{k,1:W} \sim \mathcal{N}(y_k; C_k \theta_{k,1:W}, R)$$

Observation



## Visual Tracking

(Video1) (Video2) (Video3)



## Visual Tracking – Multimodal Posteriors




The Kalman Filter loses track due to occlusion

## Visual Tracking – Multimodal Posteriors




Particle Filter with poorly designed proposal



Particle Filter with better proposal


## Visual Tracking – Multimodal Posteriors



Mixture Kalman Filter

## Sequential Monte Carlo - Particle Filtering

- We try to approximate the so-called filtering density with a set of points/Gaussians  $\equiv$  particles
- Algorithms are intuitively similar to randomised search algorithms but are best understood in terms of sequential importance sampling and resampling techniques



## Importance Sampling (IS)

Consider a probability distribution with (possibly unknown) normalisation constant

$$p(\mathbf{x}) = \frac{1}{Z} \phi(\mathbf{x}) \quad Z = \int d\mathbf{x} \phi(\mathbf{x}).$$

IS: Estimate expectations (or features) of  $p(\mathbf{x})$  by a weighted sample

$$\langle f(\mathbf{x}) \rangle_{p(\mathbf{x})} = \int d\mathbf{x} f(\mathbf{x}) p(\mathbf{x})$$

$$\langle f(\mathbf{x}) \rangle_{p(\mathbf{x})} \approx \sum_{i=1}^N \tilde{w}^{(i)} f(\mathbf{x}^{(i)})$$

## Importance Sampling (cont.)

- Change of measure with **weight function**  $W(\mathbf{x}) \equiv \phi(\mathbf{x})/q(\mathbf{x})$

$$\langle f(\mathbf{x}) \rangle_{p(\mathbf{x})} = \frac{1}{Z} \int d\mathbf{x} f(\mathbf{x}) \frac{\phi(\mathbf{x})}{q(\mathbf{x})} q(\mathbf{x}) = \frac{1}{Z} \left\langle f(\mathbf{x}) \frac{\phi(\mathbf{x})}{q(\mathbf{x})} \right\rangle_{q(\mathbf{x})} \equiv \frac{1}{Z} \langle f(\mathbf{x}) W(\mathbf{x}) \rangle_{q(\mathbf{x})}$$

- If  $Z$  is unknown, as is often the case in Bayesian inference

$$Z = \int d\mathbf{x} \phi(\mathbf{x}) = \int d\mathbf{x} \frac{\phi(\mathbf{x})}{q(\mathbf{x})} q(\mathbf{x}) = \langle W(\mathbf{x}) \rangle_{q(\mathbf{x})}$$

$$\langle f(\mathbf{x}) \rangle_{p(\mathbf{x})} = \frac{\langle f(\mathbf{x}) W(\mathbf{x}) \rangle_{q(\mathbf{x})}}{\langle W(\mathbf{x}) \rangle_{q(\mathbf{x})}}$$

## Importance Sampling (cont.)

- Draw  $i = 1, \dots, N$  independent samples from  $q$

$$\mathbf{x}^{(i)} \sim q(\mathbf{x})$$

- We calculate the **importance weights**

$$W^{(i)} = W(\mathbf{x}^{(i)}) = \phi(\mathbf{x}^{(i)})/q(\mathbf{x}^{(i)})$$

- Approximate the normalizing constant

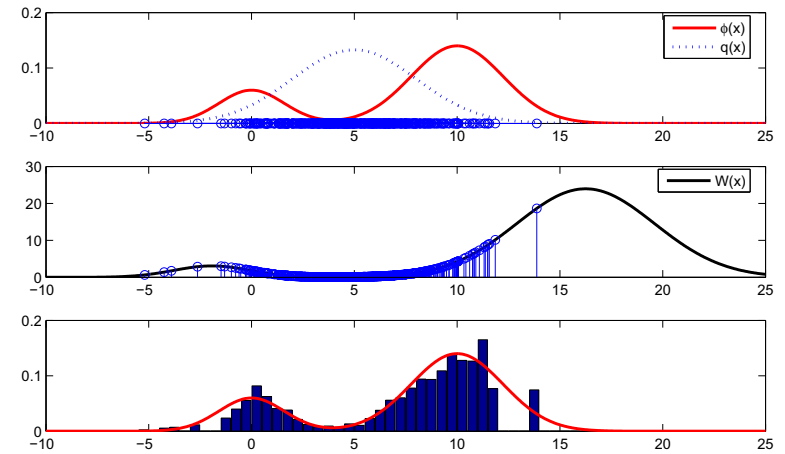
$$Z = \langle W(\mathbf{x}) \rangle_{q(\mathbf{x})} \approx \sum_{i=1}^N W^{(i)}$$

- Desired expectation is approximated by

$$\langle f(\mathbf{x}) \rangle_{p(\mathbf{x})} = \frac{\langle f(\mathbf{x}) W(\mathbf{x}) \rangle_{q(\mathbf{x})}}{\langle W(\mathbf{x}) \rangle_{q(\mathbf{x})}} \approx \frac{\sum_{i=1}^N W^{(i)} f(\mathbf{x}^{(i)})}{\sum_{i=1}^N W^{(i)}} \equiv \sum_{i=1}^N \tilde{w}^{(i)} f(\mathbf{x}^{(i)})$$

Here  $\tilde{w}^{(i)} = W^{(i)} / \sum_{j=1}^N W^{(j)}$  are *normalized importance weights*.

## Importance Sampling (cont.)



## Resampling

- Importance sampling computes an approximation with weighted delta functions

$$p(x) \approx \sum_i \tilde{W}^{(i)} \delta(x - x^{(i)})$$

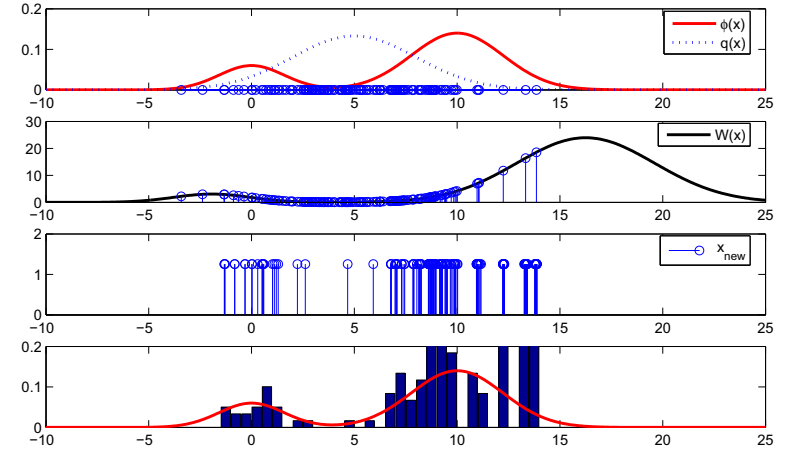
- In this representation, most of  $\tilde{W}^{(i)}$  will be very close to zero and the representation may be dominated by few large weights.
- Resampling samples a set of new “particles”

$$x_{\text{new}}^{(j)} \sim \sum_i \tilde{W}^{(i)} \delta(x - x^{(i)})$$

$$p(x) \approx \frac{1}{N} \sum_j \delta(x - x_{\text{new}}^{(j)})$$

- Since we sample from a degenerate distribution, particle locations stay unchanged. We merely duplicate (, triplicate, ...) or discard particles according to their weight.
- This process is also named “selection”, “survival of the fittest”, e.t.c., in various fields (Genetic algorithms, AI..).

## Resampling



$$x_{\text{new}}^{(j)} \sim \sum_i \tilde{W}^{(i)} \delta(x - x^{(i)})$$

## Sequential Importance Sampling, Particle Filtering

Apply importance sampling to the SSM to obtain some samples from the posterior  $p(x_{0:K}|y_{1:K})$ .

$$p(x_{0:K}|y_{1:K}) = \frac{1}{p(y_{1:K})} p(y_{1:K}|x_{0:K}) p(x_{0:K}) \equiv \frac{1}{Z_y} \phi(x_{0:K}) \quad (3)$$

Key idea: sequential construction of the proposal distribution  $q$ , possibly using the available observations  $y_{1:k}$ , i.e.

$$q(x_{0:K}|y_{1:K}) = q(x_0) \prod_{k=1}^K q(x_k|x_{1:k-1}y_{1:k})$$

## Markov Random Fields

## Markov Random Fields


- A set of random variables  $\xi = \{\xi_i\}_{i \in \mathcal{V}}$ , Given
  - an undirected graph with vertex set  $\mathcal{V}$  and undirected edge set  $\mathcal{E}$
  - A set of local potential functions (with parameters  $\mathbf{a}$ )

$$p(\xi; \mathbf{a}) = \frac{1}{Z_{\mathbf{a}}} \prod_{i \in \mathcal{V}} \phi_i(\xi_i) \prod_{(i,j) \in \mathcal{E}} \psi_{i,j}(\xi_i, \xi_j)$$

$\phi_i(\xi_i; \mathbf{a}) :$  (Singleton)

$\psi_{i,j}(\xi_i, \xi_j; \mathbf{a}) :$  (Pairwise)

## VB or Gibbs




- VB

$$q^{(\tau)}(v_k) \leftarrow \exp(\phi_k + \langle \log \psi_{k,k} + \log \psi_{k,k+1} \rangle_{q^{(\tau)}(z_k) q^{(\tau)}(z_{k+1})})$$

- Gibbs

$$v_k^{(\tau)} \sim p(v_k | z_{k-1}, z_k, y_k) \propto p(y_k | v_k) \psi_{k,k}(z_k^{(\tau)}) \psi_{k,k+1}(z_{k+1}^{(\tau)})$$

## VB or Gibbs



- VB


$$q^{(\tau)}(z_k) \leftarrow \exp(\phi_k + \langle \log \psi_{k,k-1} + \log \psi_{k,k} \rangle_{q^{(\tau)}(v_k) q^{(\tau)}(v_{k+1})})$$

- Gibbs


$$z_k^{(\tau)} \sim p(z_k | v_{k-1}, v_k) \propto \psi_{k,k-1}(v_{k-1}^{(\tau)}) \psi_{k,k}(v_k^{(\tau)})$$

## Harmonic-Transient Decomposition


- Source 1: Horizontal : Tie across time : harmonic continuity



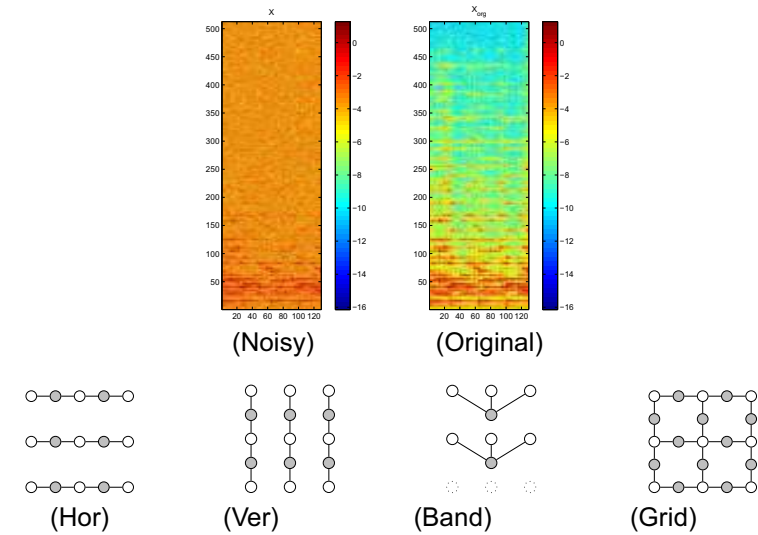
- Source 2: Vertical : Tie across frequency : transients, pulse like sounds



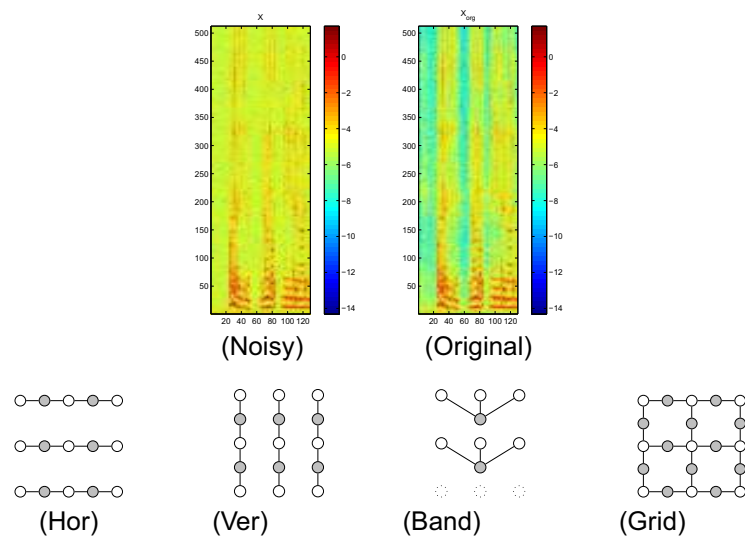
## Harmonic-Transient Decomposition



## Denoising - Piano



## Denoising - Speech



## Topic-Term-Document Models

## Text Processing, Latent Semantic Indexing

Deerwester et al. (1990), Berry et al. (1995), Manning, Schuetze, Raghavan (2007)

- We are given a database of *documents*  $D = \{d_1, \dots, d_j, \dots, d_N\}$
- Each document contains several terms from a codebook of terms  $T = \{t_1, \dots, t_i, \dots, t_M\}$
- Retrieval,
  - Given a query  $q$  (for example a set of few terms  $T_q \subset T$ ) retrieve a set of documents  $D_{\text{Retrieved}}^q$
  - Assume we know the set of relevant documents  $D_{\text{Relevant}}^q \subset D$  (with respect to the query  $q$ )
  - Quality Measures:


$$\text{Precision}^q = \frac{|D_{\text{Relevant}}^q \cap D_{\text{Retrieved}}^q|}{|D_{\text{Retrieved}}^q|} \quad \text{Recall}^q = \frac{|D_{\text{Relevant}}^q \cap D_{\text{Retrieved}}^q|}{|D_{\text{Relevant}}^q|}$$

## Representation: Term-Document matrix $A \in \mathbb{R}^{M \times N}$

- Rows : terms  $t_i, i = 1 \dots M$
- Columns: documents  $d_j = 1 \dots N$


	j.caesar	hamlet	othello	macbeth	rom&jul	sonnets
caesar	270	2	1	1	0	0
brutus	379	1	0	0	0	0
malcolm	0	0	0	60	0	0
muse	0	0	1	1	0	16
:						
love	34	68	80	19	150	195
friend	23	14	18	5	13	16
the	610	1148	759	733	682	446
traitor	1	0	0	5	1	0
traitors	9	0	1	3	0	0
:						
napkin	0	1	3	0	0	0
sword	15	16	10	14	8	1
laptop	0	0	0	0	0	0

## Term-Document matrix



- Counts

## Term-Document matrix



- Incidence (zero-one) matrix

## Singular Value Decomposition (SVD)

For any  $A \in \mathbb{R}^{M \times N}$ , there exist **orthogonal** matrices  $U \in \mathbb{R}^{M \times M}$  and  $V \in \mathbb{R}^{N \times N}$  such that

$$U = [u_1, \dots, u_M] \quad V = [v_1, \dots, v_N]$$

such that

$$U^T A V = \text{diag}(\sigma_1, \dots, \sigma_p) \in \mathbb{R}^{M \times N}$$

with  $p = \min\{M, N\}$ . We have

$$\sigma_1 \geq \sigma_2 \geq \dots \geq \sigma_p \geq 0$$

## Singular Value Decomposition (SVD)

$$A = U \times \Sigma \times V^T$$

$$\begin{pmatrix} \bullet & \bullet & \bullet & \bullet & \bullet \\ \bullet & \bullet & \bullet & \bullet & \bullet \\ \bullet & \bullet & \bullet & \bullet & \bullet \\ \bullet & \bullet & \bullet & \bullet & \bullet \end{pmatrix} = \begin{pmatrix} \bullet & \bullet & \bullet \\ \bullet & \bullet & \bullet \\ \bullet & \bullet & \bullet \end{pmatrix} \times \begin{pmatrix} \bullet & & \\ & \bullet & \\ & & \bullet \end{pmatrix} \times \begin{pmatrix} \bullet & \bullet & \bullet & \bullet & \bullet \\ \bullet & \bullet & \bullet & \bullet & \bullet \\ \bullet & \bullet & \bullet & \bullet & \bullet \\ \bullet & \bullet & \bullet & \bullet & \bullet \\ \bullet & \bullet & \bullet & \bullet & \bullet \end{pmatrix}$$

$$\begin{pmatrix} \bullet & \bullet & \bullet \\ \bullet & \bullet & \bullet \\ \bullet & \bullet & \bullet \\ \bullet & \bullet & \bullet \end{pmatrix} = \begin{pmatrix} \bullet & \bullet & \bullet & \bullet & \bullet \\ \bullet & \bullet & \bullet & \bullet & \bullet \\ \bullet & \bullet & \bullet & \bullet & \bullet \\ \bullet & \bullet & \bullet & \bullet & \bullet \end{pmatrix} \times \begin{pmatrix} \bullet & & \\ & \bullet & \\ & & \bullet \end{pmatrix} \times \begin{pmatrix} \bullet & \bullet & \bullet \\ \bullet & \bullet & \bullet \\ \bullet & \bullet & \bullet \end{pmatrix}$$

## Singular Value Decomposition (SVD)

```
>> A =
     1     2
     3     5
     1     0

>> [U S V] = svd(A)           % A == U*S*V'
U =
   -0.3559   -0.2000   -0.9129
   -0.9309   -0.0102    0.3651
   -0.0823    0.9797   -0.1826

S =
    6.2638         0
         0    0.8744
         0         0

V =
   -0.5158    0.8567
   -0.8567   -0.5158
```

## Singular Value Decomposition (SVD)

```
>> U(:,1)*S(1,1)*V(:,1)'
ans =
    1.1498    1.9098
    3.0076    4.9954
    0.2661    0.4419

>> U(:,2)*S(2,2)*V(:,2)'
ans =
   -0.1498    0.0902
   -0.0076    0.0046
    0.7339   -0.4419

>> U(:,1)*S(1,1)*V(:,1)' + U(:,2)*S(2,2)*V(:,2)'   %% == U*S*V' == A
ans =
    1.0000    2.0000
    3.0000    5.0000
    1.0000    0.0000

>> A =
     1     2
     3     5
     1     0
```

## Singular Value Decomposition (SVD)

SVD expansion


$$\begin{aligned} A &= \sum_{r=1}^P \sigma_r u_r v_r^\top \\ &= U \mathbf{diag}(\sigma_1, \dots, \sigma_P) V^\top \end{aligned}$$

The norm relations for  $A \in \mathbb{R}^{M \times N}$ ,  $P = \min\{M, N\}$

$$\begin{aligned} \|A\|_F^2 &= \sigma_1^2 + \dots + \sigma_P^2 \\ \|A\|_2^2 &= \sigma_1^2 \end{aligned}$$


## Singular Value Decomposition of Term-Document Matrices

Another “term-document” matrix




## Singular Value Decomposition of Term-Document Matrices

$$A \approx U(:, 1:n) S(1:n, 1:n) V(:, 1:n)^\top$$




## Singular Value Decomposition of Term-Document Matrices

$$A \approx U(:, 1:n) S(1:n, 1:n) V(:, 1:n)^\top$$






## Rank-1 Matrices



## Rank-1 Matrices



## LSI: Summary and Remarks


- Low rank approximation to a term-document matrix by keeping the latent dimensions corresponding to  $n$  largest singular values of SVD

$$A \approx \sum_{r=1}^n \sigma_r U(:,r) V(:,r)^\top$$

- No direct statistical interpretation, but loosely
  - Each  $r = 1 \dots n$  denotes a *latent topic* ( $n$  is the total number of topics)
  - $U(i, r)$  corresponds to *weight* of the  $i$ 'th term given the topic  $r$
  - $V(j, r)$  corresponds to *emphasis* of topic  $r$  in document  $j$   
 We can think  $V(j, 1:n)^\top$  as the coordinates of  $j$ 'th document in an  $n$  dimensional *latent topic space*
  - The coordinates of a new document are computed by

$$v_{\text{new}} = \Sigma^{-1} U^\top a_{\text{new}}$$

## Latent Semantic Space



## LSI: Summary and Remarks

- Clustering, assessing similarity, visualisation ...
- Rationale: documents that share frequent co-occurring terms will be close in the latent space
- May deal with synonymy and polysemy
  - different words - same meaning  
baggage-lugagge
  - same word - different meaning  
spider (the animal - the web crawler)

## Probabilistic Latent Sematic Indexing, the Aspect Model

Hofmann, 1999

$d \sim p(d)$	Document
$z d \sim p(z d)$	Latent Topic
$t z \sim p(t z)$	Term

More to come ...

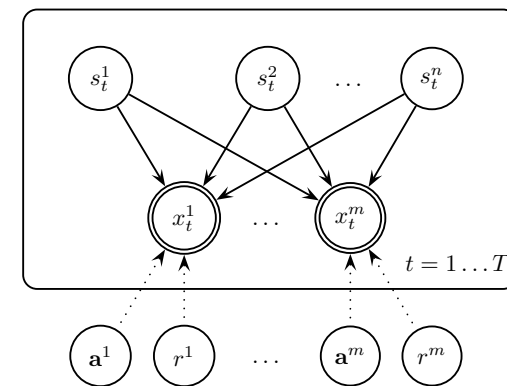
## Factorial Models

### Source Separation

### Bayesian Model selection

## Audio Source Separation

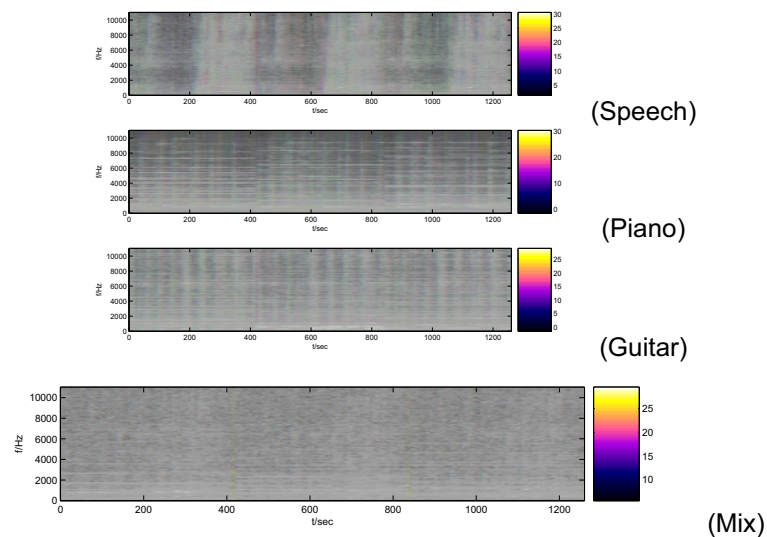
Estimate  $n$  hidden signals  $s_t$  from  $m$  observed signals  $x_t$ .



$$s_t^i \sim p(s_t^i)$$

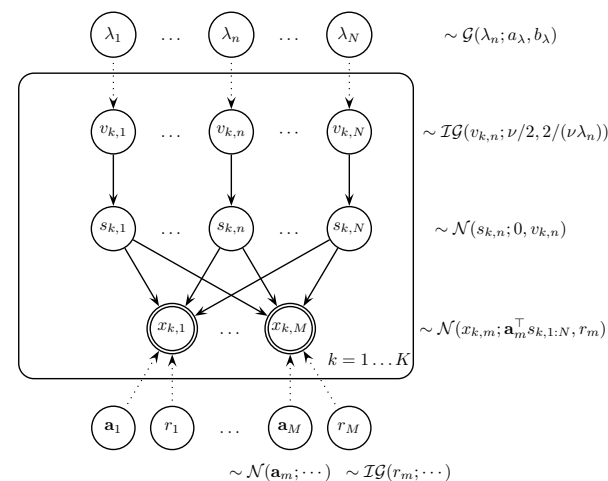
$$x_t^j \sim \mathcal{N}(x; \mathbf{a}^j s_t^{1:n}, r^j)$$

## Audio Source Separation

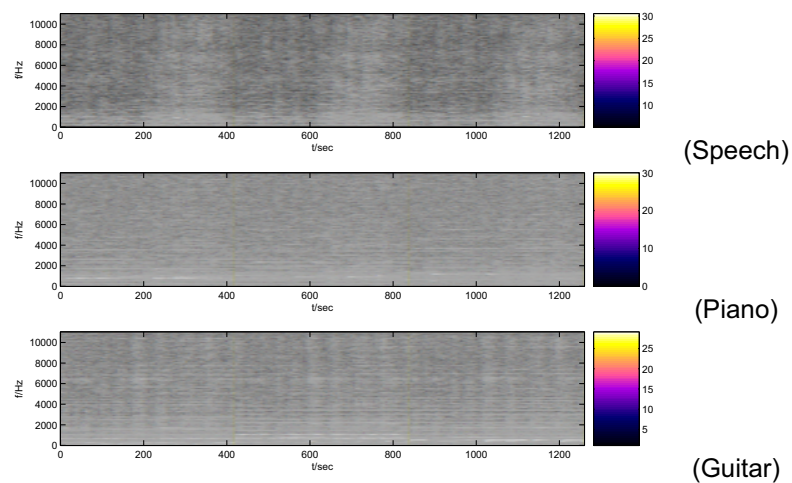


## Audio Source Separation


- Hierarchical Prior Model (Fevotte and Godsill 2005 [10], Cemgil et. al. 2006 [3])



## Reconstructions




## Audio Source Separation, Inference




- Exact inference is not possible

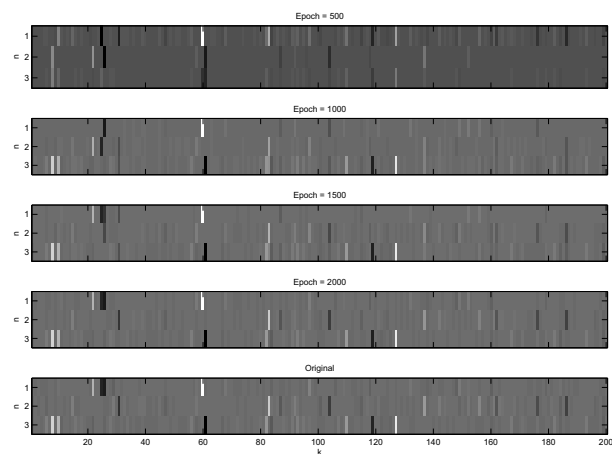
## Observations



## A typical run, 250/250 Gibbs/VB with tempering




## Reconstructions



Posterior surface is multimodal, each mode corresponding to a viable separation


## Bayesian Variable Selection



- Generalized Linear Model – Column's of  $C$  are the basis vectors
- The exact posterior is a mixture of  $2^W$  Gaussians
- When  $W$  is large, computation of posterior features becomes intractable.

## Generative model


$$\begin{aligned}
 r_i &\sim \mathcal{C}(r_i; \pi) \\
 s_i | r_i &\sim \mathcal{N}(s_i; \mu(r_i), \Sigma(r_i)) \\
 \mathbf{x} | s_{1:W} &\sim \mathcal{N}(\mathbf{x}; C s_{1:W}, R) \\
 C &\equiv [ C_1 \quad \dots \quad C_i \quad \dots \quad C_W ]
 \end{aligned}$$



$$p(\mathbf{x}, s_{1:W}, r_{1:W}) = p(\mathbf{x} | s_{1:W}, r_{1:W}) \prod_{i=1}^W p(s_i | r_i) p(r_i)$$

## Example 1: Variable selection in Polynomial Regression

Given  $\{t_j, x(t_j)\}_{j=1 \dots J}$ , what is the order  $N$  of the polynomial?



$$x(t) = \sum_{i=0}^N s_{i+1} t^i + \epsilon(t)$$

## Ex1: Regression

$$\begin{aligned}
 \mathbf{t} &= (t_1 \quad t_2 \quad \dots \quad t_J)^\top \\
 C &\equiv (\mathbf{t}^0 \quad \mathbf{t}^1 \quad \dots \quad \mathbf{t}^{W-1})
 \end{aligned}$$

```
>> C = fliplr(vander(0:4)) % Van der Monde matrix
    1     0     0     0     0
    1     1     1     1     1
    1     2     4     8    16
    1     3     9    27    81
    1     4    16    64   256
```

$$\begin{aligned}
 r_i &\sim \mathcal{C}(r_i; 0.5, 0.5) \quad r_i \in \{\text{on}, \text{off}\} \\
 s_i | r_i &\sim \mathcal{N}(s_i; 0, \Sigma(r_i)) \\
 \mathbf{x} | s_{1:W} &\sim \mathcal{N}(\mathbf{x}; C s_{1:W}, R)
 \end{aligned}$$

$$\Sigma(r_i = \text{on}) \gg \Sigma(r_i = \text{off})$$

## Ex1: Regression


To find the “active” basis functions we need to calculate

$$r_{1:W}^* \equiv \underset{r_{1:W}}{\operatorname{argmax}} p(r_{1:W} | \mathbf{x}) = \underset{r_{1:W}}{\operatorname{argmax}} \int ds_{1:W} p(\mathbf{x} | s_{1:W}) p(s_{1:W} | r_{1:W}) p(r_{1:W})$$


Then, the reconstruction is given by

$$\begin{aligned}
 \hat{x}(t) &= \left\langle \sum_{i=0}^{W-1} s_{i+1} t^i \right\rangle_{p(s_{1:W} | \mathbf{x}, r_{1:W}^*)} \\
 &= \sum_{i=0}^{W-1} \langle s_{i+1} \rangle_{p(s_{i+1} | \mathbf{x}, r_{1:W}^*)} t^i
 \end{aligned}$$


## Ex1: Regression



## Ex1: Regression



## Example 2: Chord Recognition



## (Damped) Sinusoidal Basis

- $h = 1 \dots H$ , number of harmonics,  $t = 0 \dots T - 1$ , sample index
- $\omega$  : fundamental frequency in rad,  $\rho$  damping coefficient

$$C(\omega) \equiv \begin{pmatrix} C_0^1 & \dots & C_0^H \\ \vdots & C_t^h & \vdots \\ C_{T-1}^1 & \dots & C_{T-1}^H \end{pmatrix}$$


$$C_t^h \equiv \rho^t \begin{pmatrix} \cos(th\omega) & \sin(th\omega) \end{pmatrix}$$

$$\mathbf{C} = [C(\omega_1) \dots C(\omega_\nu) \dots C(\omega_W)]$$


- See also Badeau, Boyer, David. Eds parametric modelling and tracking of audio signals. In DAFx 2002

## Factor graph

$$\begin{aligned} \log \phi(r_{1:W}, s_{1:W}) &= \sum_{i=1}^W (\log \pi(r_i)) \\ &+ \sum_{i=1}^W \left( -\frac{1}{2} s_i^\top \Sigma(r_i)^{-1} s_i + \mu(r_i)^\top \Sigma(r_i)^{-1} s_i \right. \\ &\quad \left. - \frac{1}{2} \mu(r_i)^\top \Sigma(r_i)^{-1} \mu(r_i) - \frac{1}{2} \log |2\pi \Sigma(r_i)| \right) \\ &- \frac{1}{2} \mathbf{x}^\top R^{-1} \mathbf{x} + s_{1:W}^\top C^\top R^{-1} \mathbf{x} - \frac{1}{2} s_{1:W}^\top C^\top R^{-1} C s_{1:W} - \frac{1}{2} \log |2\pi R| \end{aligned}$$




## Approximating Structures



$$Q_1 = \prod_{i=1}^W Q(s_i) Q(r_i) \quad Q_2 = Q(s_{1:W}) \prod_{i=1}^W Q(r_i) \quad Q_3 = \prod_{i=1}^W Q(s_i, r_i)$$

## MCMC versus Variational Bayes (VB)

- Each configuration of  $r_{1:W}$  corresponds to a corner of a  $W$  dimensional hypercube

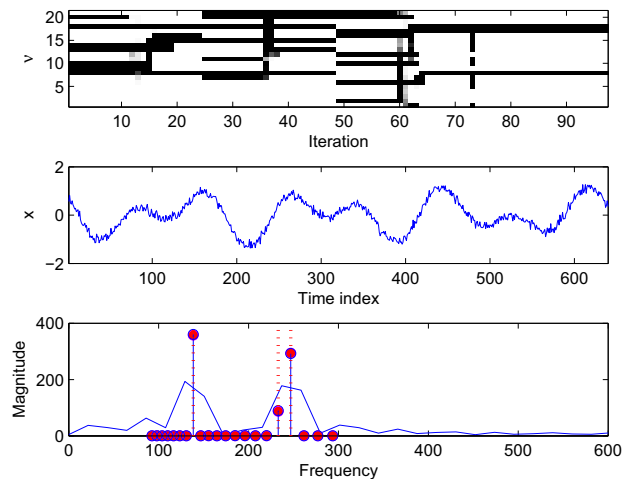


- MCMC** moves along the edges stochastically
- Iterative Improvement** moves along the edges greedily
- VB** moves inside the hypercube deterministically

## Iterative Improvement


iteration	$r_1$	$r_M$	$\log p(y_{1:T}, r_{1:M})$
1	o o o o o o o o	o o o o o o o o	-1220638254
2	o o o o o o o o	o o o o o o o o	-665073975
3	o o o o o o o o	o o o o o o o o	-311983860
4	o o o o o o o o	o o o o o o o o	-162334351
5	o o o o o o o o	o o o o o o o o	-43419569
6	o o o o o o o o	o o o o o o o o	-1633593
7	o o o o o o o o	o o o o o o o o	-14336
8	o o o o o o o o	o o o o o o o o	-5766
9	o o o o o o o o	o o o o o o o o	-5210
10	o o o o o o o o	o o o o o o o o	-4664
True	o o o o o o o o	o o o o o o o o	-4664

## Results, VB with tempering and reinitialisation



$F_s = 22050$  Hz,  $N = 29$  msec,  $H = 1$ , Midinotes = 30...50

## Results, MCMC with tempering and reinitialisation



$F_s = 22050$  Hz,  $N = 29$  msec,  $H = 1$ , Midinotes = 30...50

## Bayesian/Generative/Probabilistic approaches to Polyphonic Transcription

(Walmsley 2000, Davy and Godsill 2002, Raphael 2001, Abdallah 2002, Cemgil et. al. 2003-2006, Vincent 2003, Vincent and Plumbley 2005, Vogel, Jordan and Wessel 2005, Thornburg, Leitsnikov and Berger 2004, Blumensath and Davies 2006, Dubois and Davy 2005)

- Various related but different models
- Inference schemata
  - Reversible Jump MCMC
  - Iterative Improvement
  - Laplace approximation
  - Particle filtering
  - Variational Bayes, MCMC

## Summary

- Bayesian Inference
- Graphical models
- Exact Inference
- Approximate inference



## Summary, Attributes of Probabilistic Inference

- **Exact** ↔ **Approximate**
- **Deterministic** ↔ **Stochastic**
- **Online** ↔ **Offline**
- **Centralized** ↔ **Distributed**

## Summary of what we have mentioned

- Exact inference, Belief Propagation
- Approximate inference
  - Deterministic
    - \* Variational Bayes,
    - \* Expectation/Maximization (EM), Iterative Conditional Modes (ICM)
  - Stochastic
    - \* Markov Chain Monte Carlo
    - \* Importance Sampling,
    - \* Particle filtering

## Summary of what we have not mentioned

- Exact Inference (Junction Tree ...)
- Deterministic Inference
  - Assumed Density Filter (ADF), Extended Kalman Filter (EKF), Unscented Particle Filter
  - Structured Mean field
  - Loopy Belief Propagation, Expectation Propagation, Generalized Belief Propagation
  - Fractional Belief propagation, Bound Propagation, <your favorite name> Propagation
  - Graph cuts ...
- Stochastic
  - Unscented Particle Filter, Nonparametric Belief Propagation
  - Annealed Importance Sampling, Adaptive Importance Sampling
  - Hybrid Monte Carlo, Exact sampling, Coupling from the past

## Bibliography

- General background about probability theory
- Graphical models
- Exact inference
- Variational Methods
- Markov Chain Monte Carlo
- Sequential Monte Carlo
- Applications

## General background about probability theory

- Dimitri P. Bertsekas and John N. Tsitsiklis. Introduction to Probability. Athena Scientific, 2002
- Geoffrey Grimmet and David Stirzaker, Probability and Random Processes, (3rd Ed), Oxford, 2006

## “Instant Classics” of Bayesian Machine Learning and Graphical Models

- Michael I. Jordan, Learning in Graphical Models, 1998
- David MacKay Information Theory, Learning and Inference Algorithms, 2003, Cambridge
- Chris Bishop, Machine Learning and Pattern Recognition, 2006, Springer

## Further Reading, Variational Methods

- Jaakkola “Tutorial on variational approximation methods”, 2000  
<http://people.csail.mit.edu/tommi/papers/Jaa-var-tutorial.ps>
- Wainwright and Jordan 2003 [19] Berkeley EECS Tech. Rep.
- Frey and Jojic, PAMI 2005 [11]
- Winn and Bishop “Variational Message Passing” 2005 JMLR [20]

## Further Reading, MCMC and SMC tutorials and overviews

- Andrieu, de Freitas, Doucet, Jordan. *An Introduction to MCMC for Machine Learning*, 2001
- Andrieu. *Monte Carlo Methods for Absolute beginners*, 2004
- Doucet, Godsill, Andrieu. “On Sequential Monte Carlo Sampling Methods for Bayesian Filtering”, Statistics and Computing, vol. 10, no. 3, pp. 197-208, 2000
- Gilks, Richardson, Spiegelhalter, *Markov Chain Monte Carlo in Practice*, Chapman Hall, 1996
- Doucet, de Freitas, Gordon, *Sequential Monte Carlo Methods in Practice*, Springer, 2001

## Text Processing and Information Retrieval

- *Information Retrieval*, Manning, Schuetze, and Raghavan, Cambridge University Press, 2007 (Draft)  
<http://www-csli.stanford.edu/~schuetze/information-retrieval-book.html>
- *Modeling the Internet and the Web: Probabilistic Methods and Algorithms*. Pierre Baldi, Paolo Frasconi, Padhraic Smyth, Wiley, 2003  
<http://ibook.ics.uci.edu>

Compression, Efficient Data Structures,

- *Managing Gigabytes: Compressing and Indexing Documents and Images*. Witten, Moffat and Bell. Morgan Kaufmann 1999
- *Introduction to Data Compression*. Khalid Sayood, Morgan Kaufmann (3rd Ed), 2005

## Some Generic Software Packages

- Kevin Murphy's Matlab Bayesian Networks toolkit (BNT)
- Gilks, et. al. BUGS, WinBUGS – (Bayesian analysis Using Gibbs Sampling) A powerful program that compiles Gibbs Samplers from
- Winn, et. al, VIBES – Similar to BUGS but for variational inference

For source separation, there are some specialised libraries

- Petersen and Winther (DTU, Copenhagen)
- Harva, Raiko, Honkela, Valpola “Bayes Blocks” (HUT, Helsinki)

## Music Applications

- Klapuri and Davy (Eds) *Signal processing for Music Transcription*, Springer, 2006
- Temperley, *Probability and Music*, MIT Press, 2007

## References

- [1] M. Allan and C. K. I. Williams. Harmonising chorales by probabilistic inference. In *Advances in Neural Information Processing Systems 17*, 2004.
- [2] J.E. Besag. On the statistical analysis of dirty pictures (with discussion). *Jr. R. Stat. Soc. B*, 48:259–302, 1986.
- [3] A. T. Cemgil, C. Févotte, and S. J. Godsill. Variational and Stochastic Inference for Bayesian Source Separation. *Digital Signal Processing*, in Print, 2007.
- [4] A. T. Cemgil and S. J. Godsill. Efficient Variational Inference for the Dynamic Harmonic Model. In *Proc. of IEEE Workshop on Applications of Signal Processing to Audio and Acoustics*, New Paltz, NY, October 2005.
- [5] A. T. Cemgil and S. J. Godsill. Probabilistic Phase Vocoder and its application to Interpolation of Missing Values in Audio Signals. In *13th European Signal Processing Conference*, Antalya/Turkey, 2005. EURASIP.
- [6] A. T. Cemgil and H. J. Kappen. Monte Carlo methods for Tempo Tracking and Rhythm Quantization. *Journal of Artificial Intelligence Research*, 18:45–81, 2003.
- [7] A. T. Cemgil, H. J. Kappen, and D. Barber. A Generative Model for Music Transcription. *IEEE Transactions on Audio, Speech and Language Processing*, 14(2):679–694, March 2006.
- [8] A.T. Cemgil, H. J. Kappen, P. Desain, and H. Honing. On tempo tracking: Tempogram Representation and Kalman filtering. In *Proceedings of the 2000 International Computer Music Conference*, pages 352–355, Berlin, 2000. (This paper has received the Swets and Zeitlinger Distinguished Paper Award of the ICMC 2000).
- [9] R. Chen and J. S. Liu. Mixture Kalman filters. *J. R. Statist. Soc.*, 10, 2000.
- [10] C. Févotte and S. J. Godsill. A Bayesian approach for blind separation of sparse sources. *IEEE Trans. Speech and Audio Processing*, in press. In press - Preprint available at <http://persos.mist-technologies.com/~cfévotte/>.

- [11] B. J. Frey and N. Jojic. A comparison of algorithms for inference and learning in probabilistic graphical models. IEEE Transactions on Pattern Analysis and Machine Intelligence, 27(9), 2005.
- [12] Z. Ghahramani and M. Beal. Propagation algorithms for variational Bayesian learning. In Neural Information Processing Systems 13, 2000.
- [13] E. T. Jaynes. Probability Theory, The Logic of Science. Cambridge University Press, edited by G. L. Bretthorst, 2003.
- [14] F. R. Kschischang, B. J. Frey, and H.-A. Loeliger. Factor graphs and the sum-product algorithm. IEEE Transactions on Information Theory, 47(2):498–519, February 2001.
- [15] D. J. C. MacKay. Information Theory, Inference and Learning Algorithms. Cambridge University Press, 2003.
- [16] Radford M. Neal and Geoffrey E. Hinton. A view of the EM algorithm that justifies incremental, sparse, and other variants. In Learning in graphical models, pages 355–368. MIT Press, 1999.
- [17] L. R. Rabiner. A tutorial in hidden Markov models and selected applications in speech recognition. Proc. of the IEEE, 77(2):257–286, 1989.
- [18] C. Raphael. A probabilistic expert system for automatic musical accompaniment. Journal of Computational and Graphical Statistics, 10(3):467–512, 2001.
- [19] M. Wainwright and M. I. Jordan. Graphical models, exponential families, and variational inference. Technical Report 649, Department of Statistics, UC Berkeley, September 2003.
- [20] J. Winn and C. Bishop. Variational message passing. Journal of Machine Learning Research, 6:661–694, 2005.

**Thank you for your patience and attention!**

Slides will be available online

<http://www-sigproc.eng.cam.ac.uk/~atc27/acm-tutorial/>